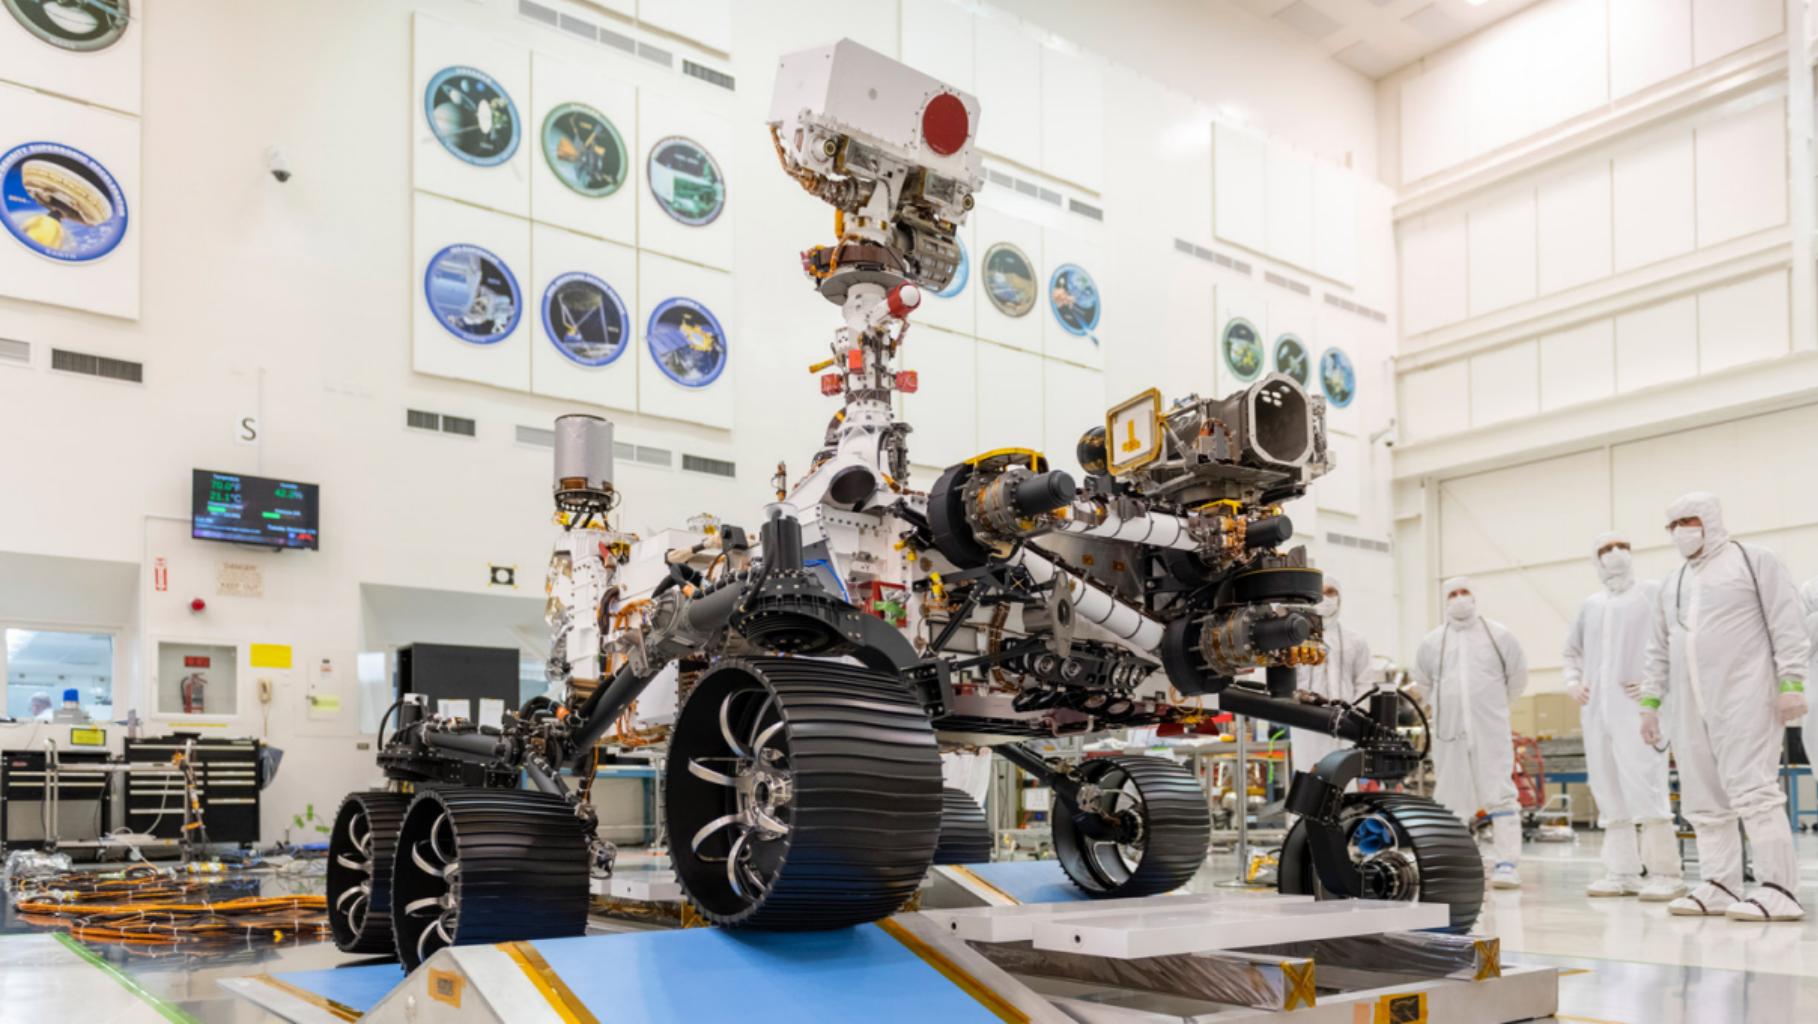




ENVIRONMENT LEARNING FOR ROBOT PATH-PLANNING

LUCAS SALDYT ARIZONA STATE UNIVERSITY







Core Goals & Concepts

Learning

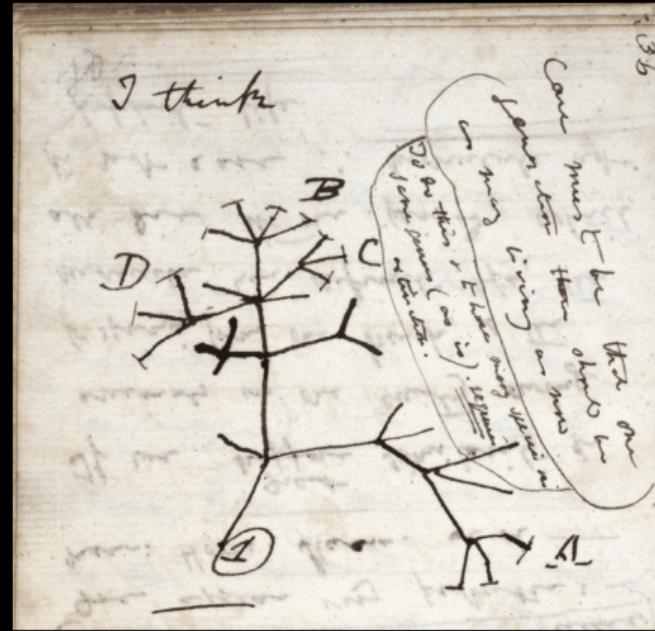
- Flexibility: adapt to new tasks

Generalization

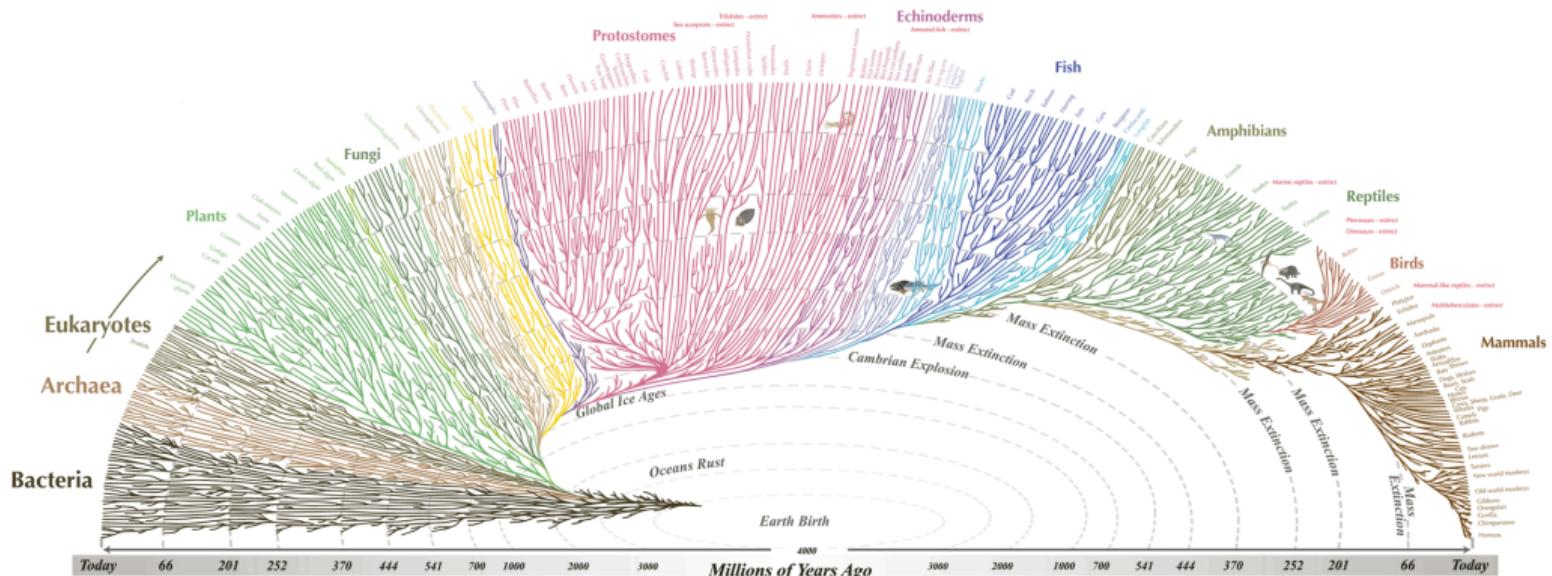
- Learn overall task proficiency?

Specialization

- Learn environment details?



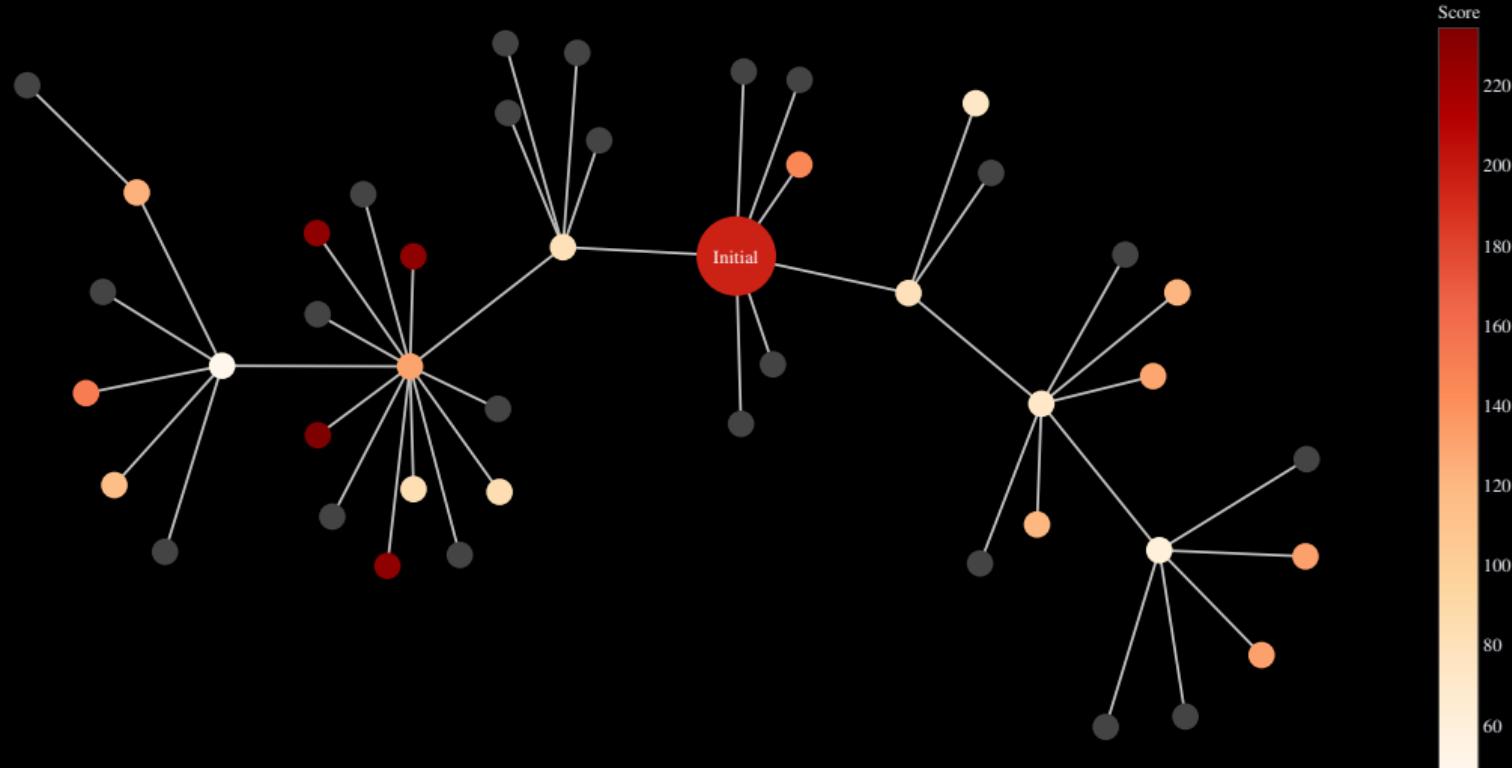
Where is this image from?



Natural evolution

Computer Programs
instead of species

Evolutionary Programming



Evolved Computer Programs

Problems Considered

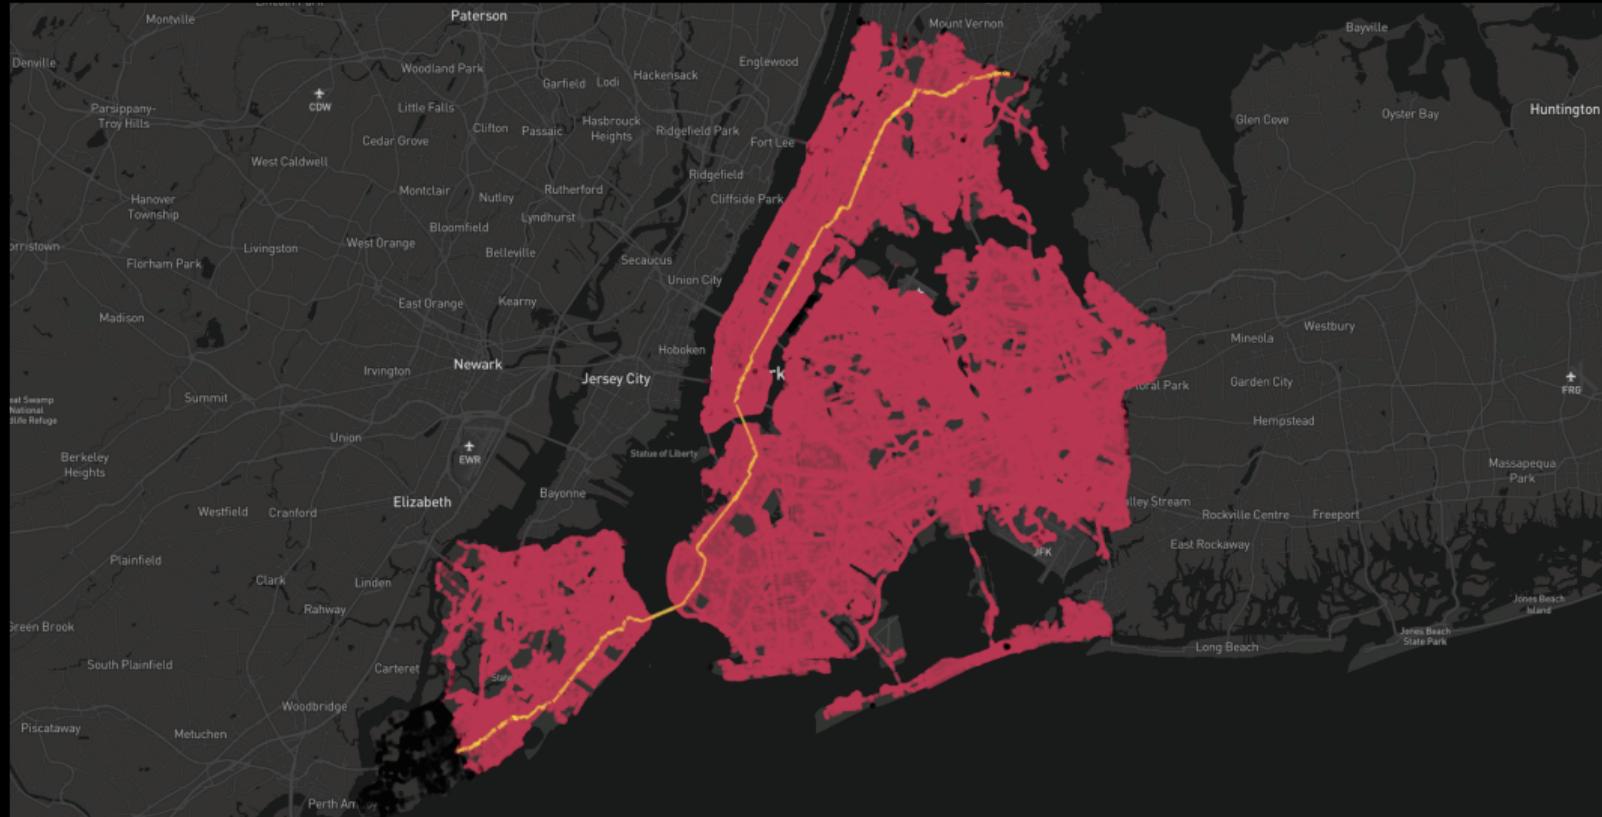


Path Planning

- Graph-based
- Sample-based (left)
- Dynamic (i.e. Mars)
- Multiagent (i.e. Amazon)

Symbolic Regression
Neural Arch. Search

Graph search w/ Dijkstra (New York)

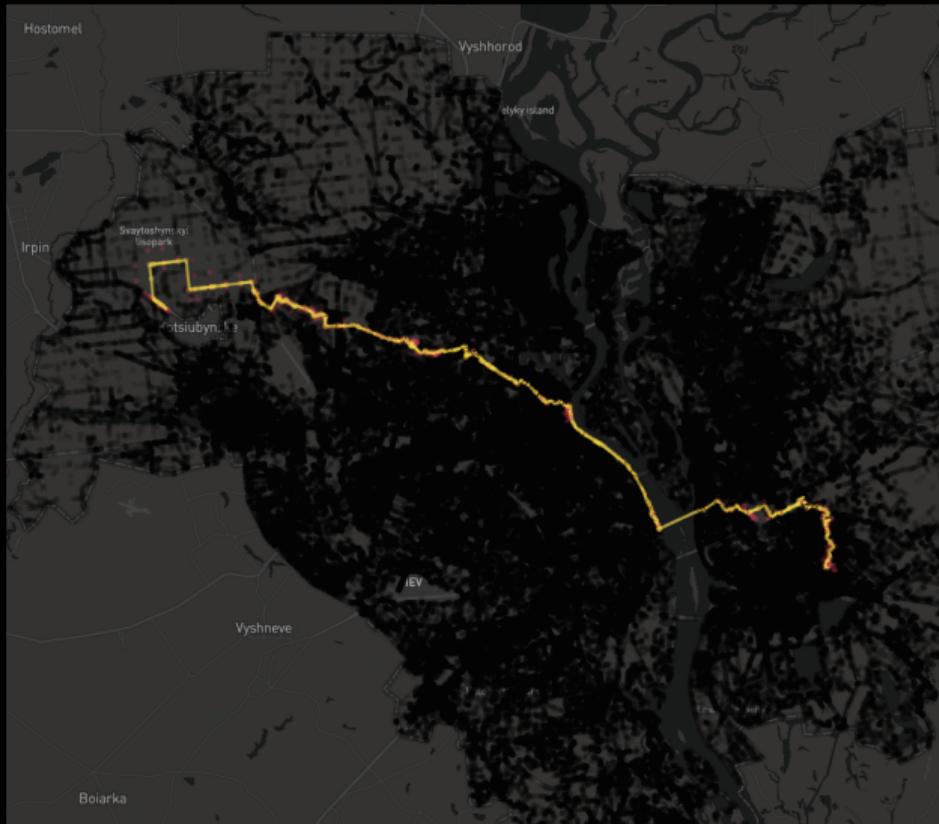


Graph search w/ A* (New York)



Graph-based Algorithms (Kiev, Ukraine)

- Breadth-first (BFS)
- Depth-first (DFS)
- Dijkstra/A*
- Goal-directed
- Separator-based
- Hierarchical
- Bounded Hop
- Hybrid Techniques
- Bast et al., 2016



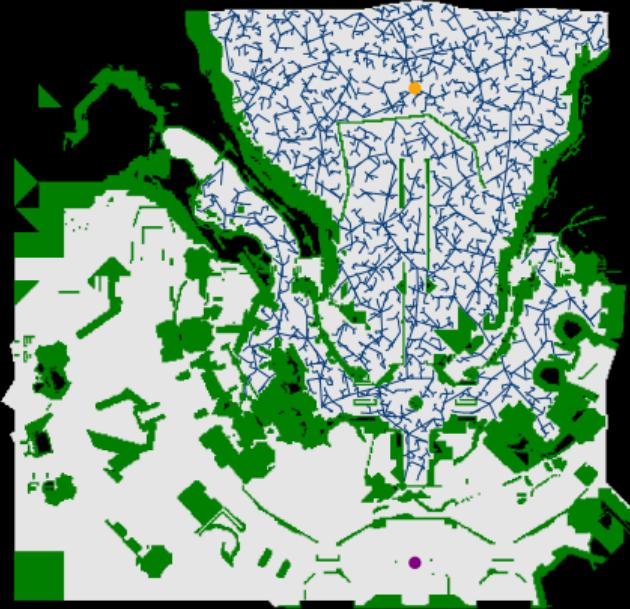
Sampling-based search

- Rapidly Exploring Random Trees (RRTs)

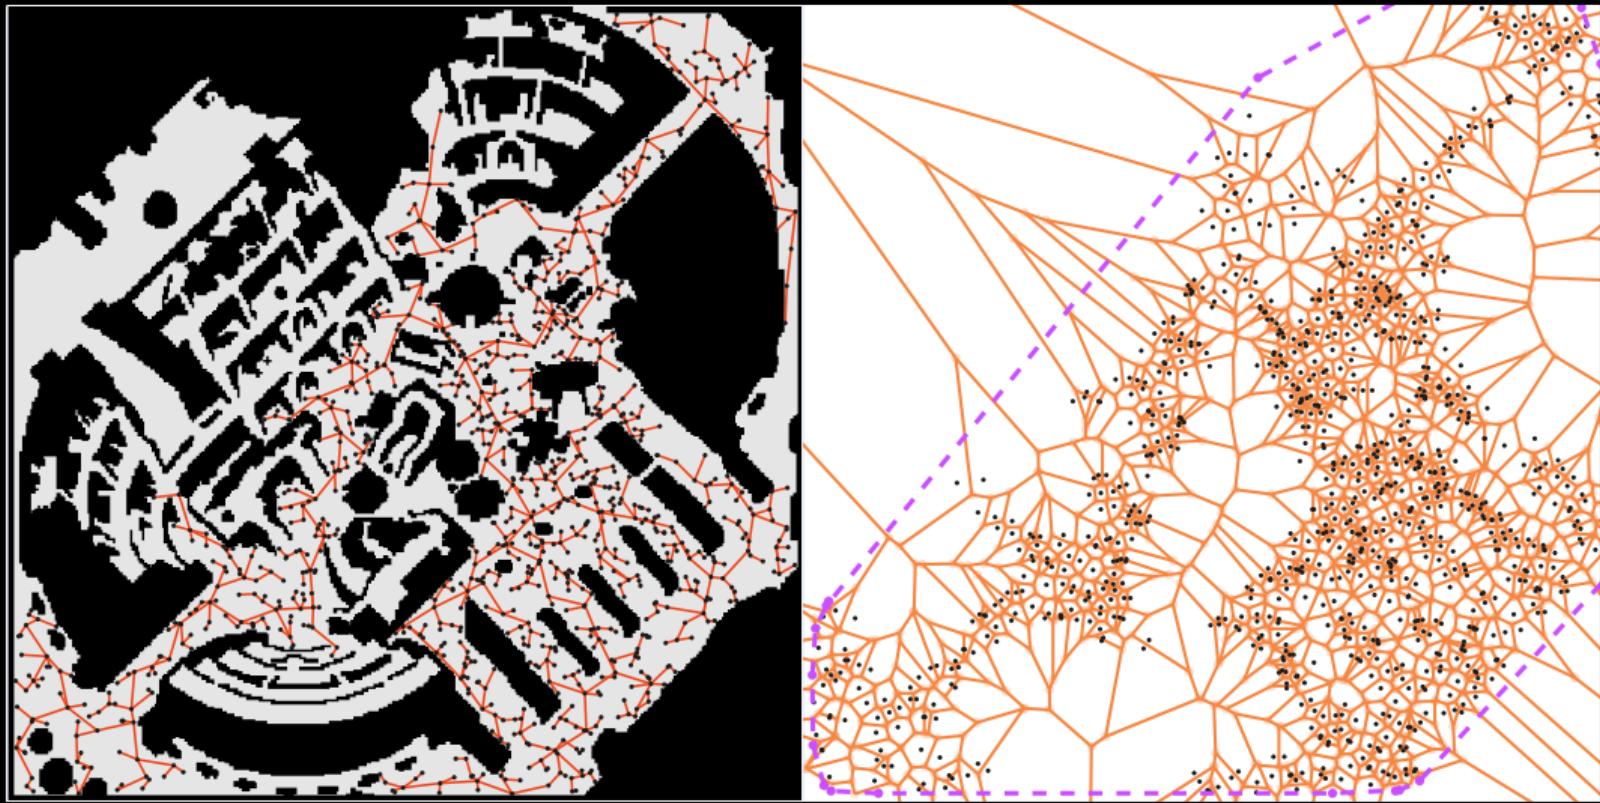
- RRT, RRT*, Informed RRT*
- Constrained (RRT-Blossom)
- Multi-phase (A*-RRT, Theta*-RRT)
- Dynamic (RT-RRT*, RRTX, RRT)
- LaValle and Kuffner Jr, 2001...

- Probabilistic Road Maps (PRMs)

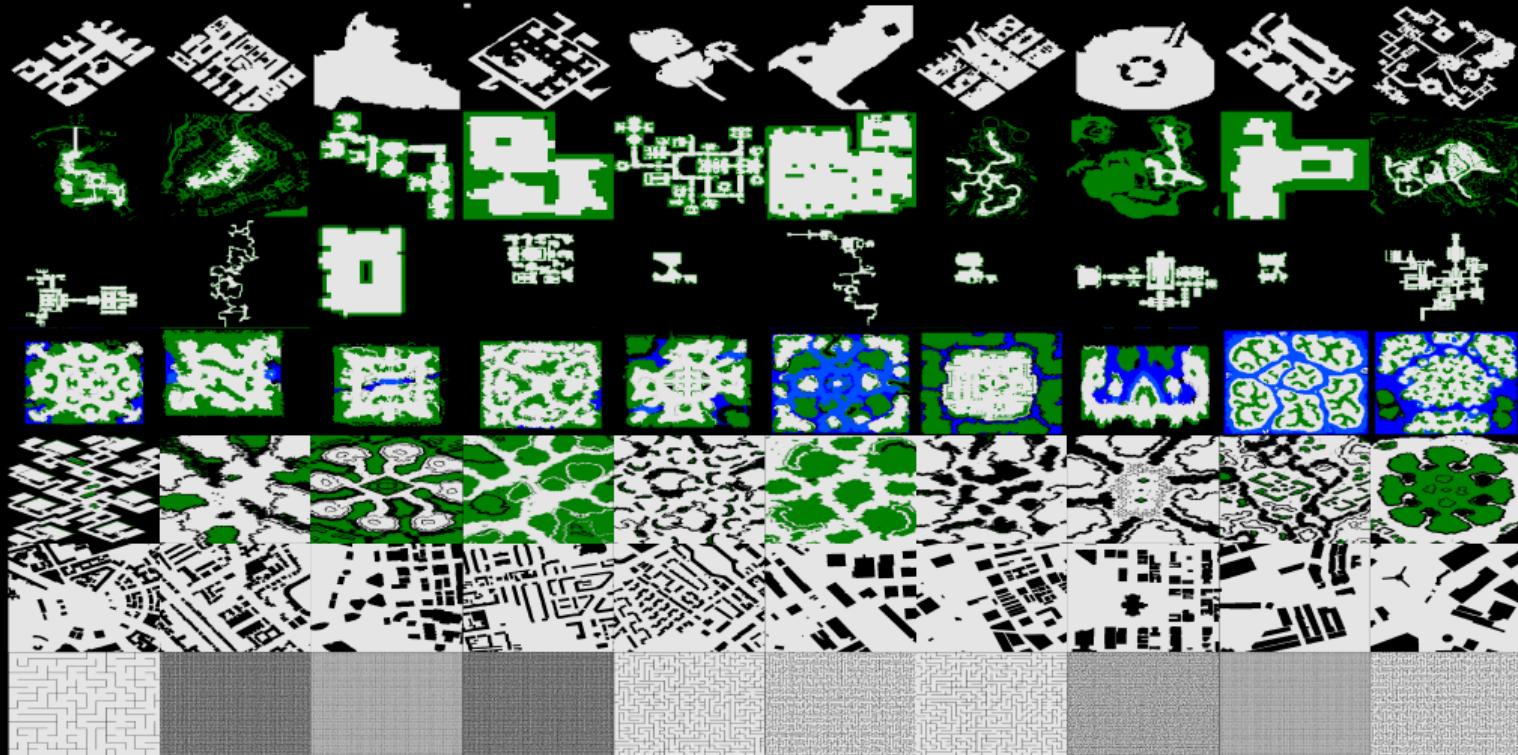
- Combine sampling & graph-based
- Kavraki et al., 1996



RRT* (Baldur's Gate)



Maps Overview (Sturtevant, 2012)



Maps Detail (Sturtevant, 2012)

Streetmaps: Boston_0_1024 Streetmaps: Shanghai_0_1024

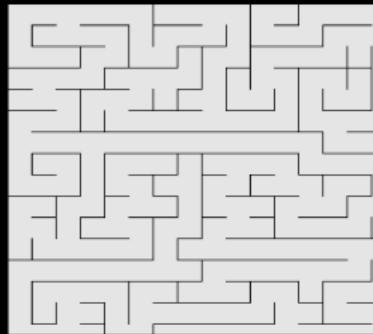


Dragon Age: orz100d



Dragon Age: ht_mansion2b

Maze: maze512-32-9

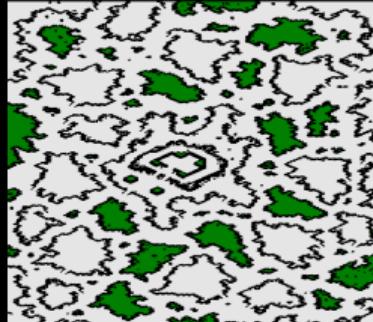


Starcraft: TheFrozenSea

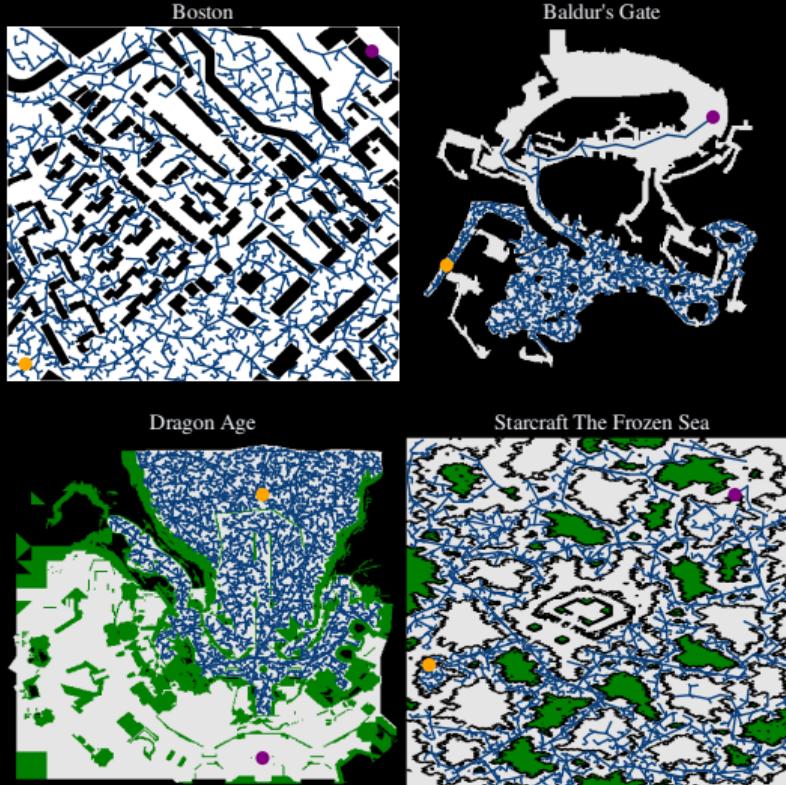
Baldurs Gate: AR0300SR



Starcraft: WheellofWar

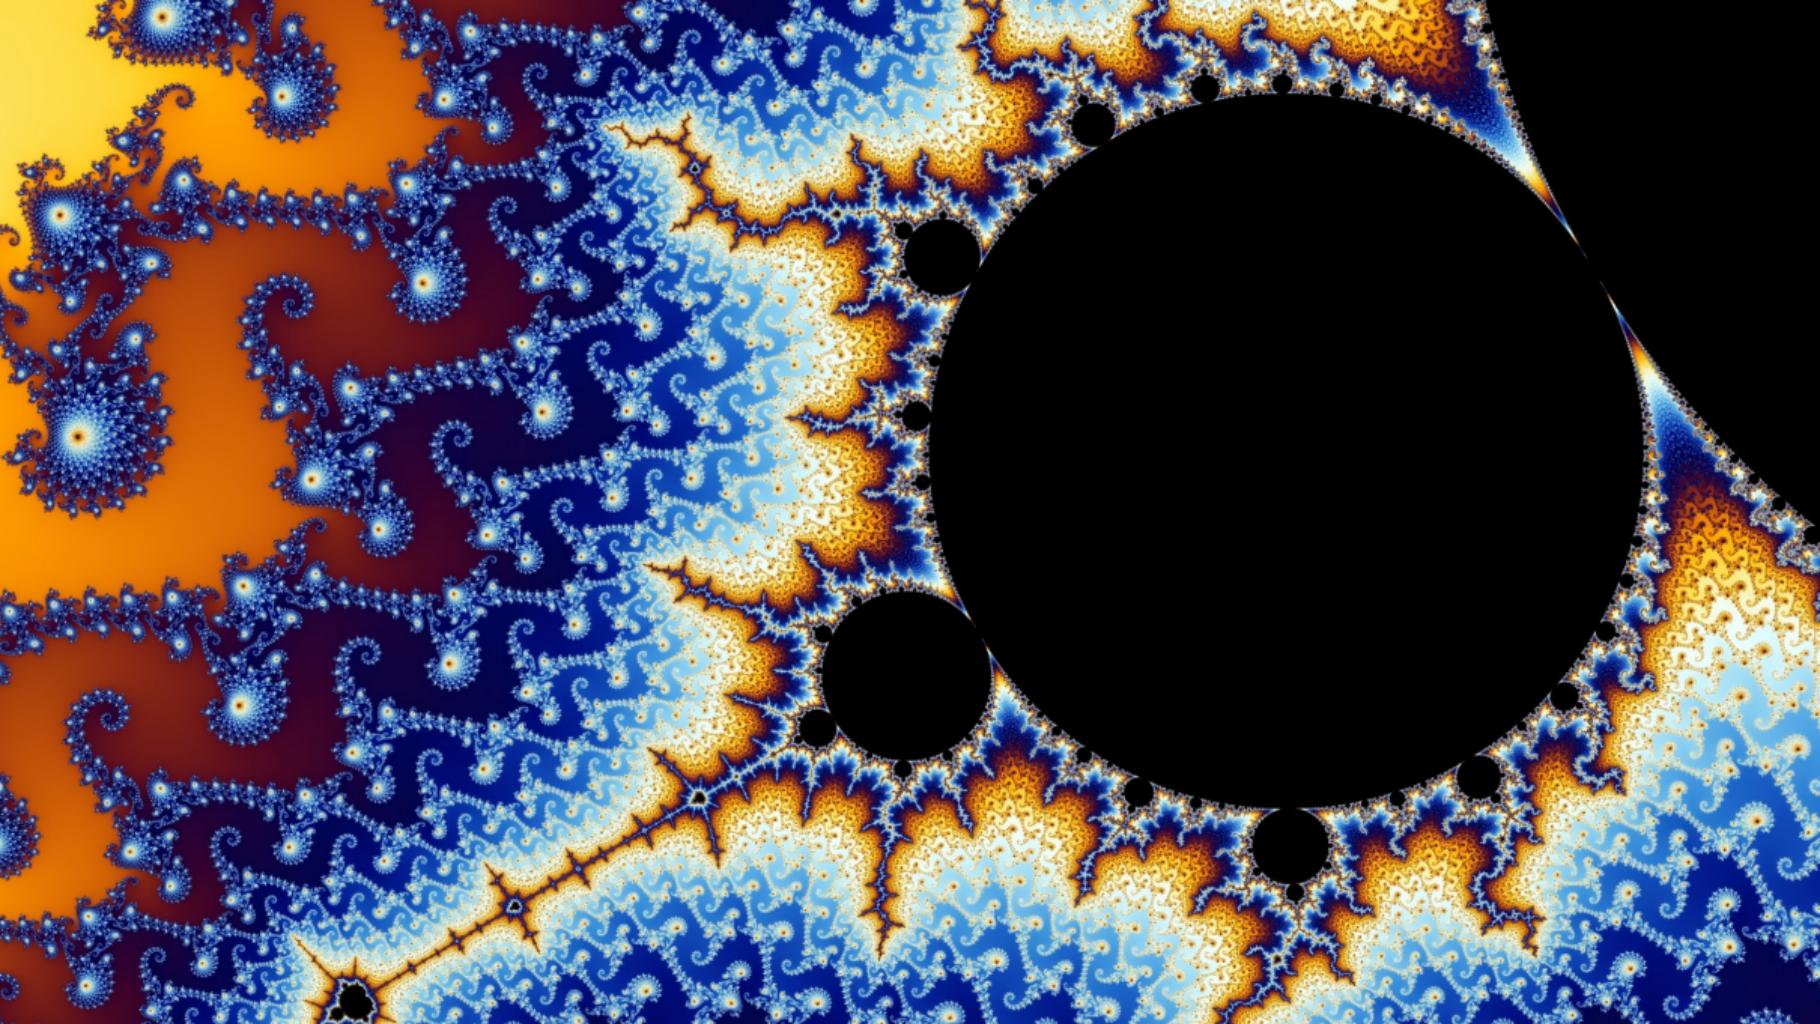


Specialization Teaser



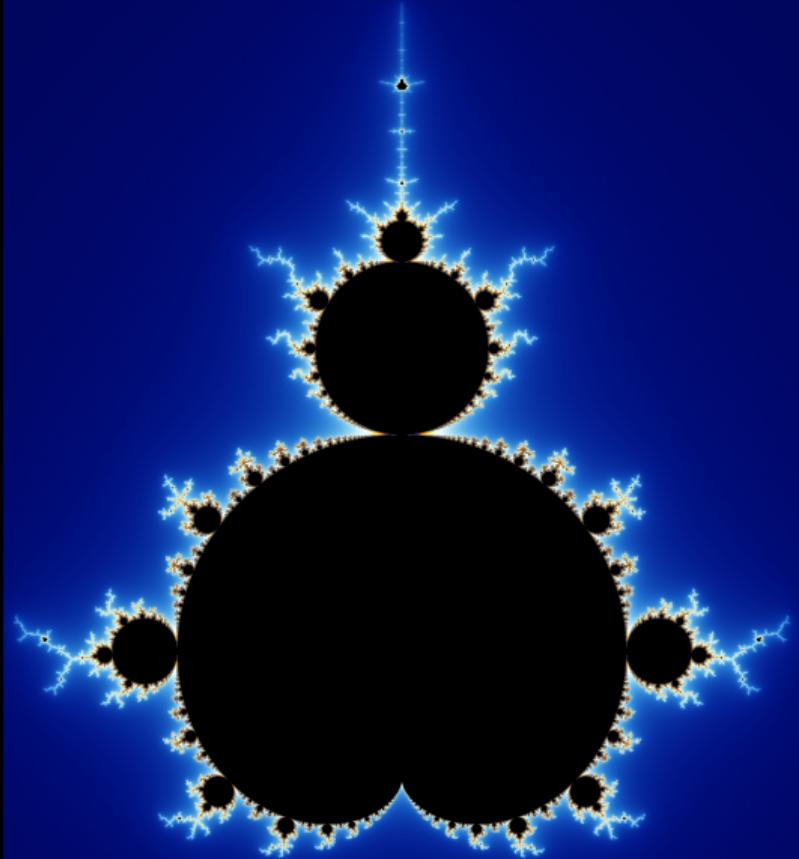
There is no “silver bullet” algorithm for solving all path planning problems. Heuristics that lead to massive speed-up in one scenario might be detrimental in others. Also, algorithmic parameters are mostly ad-hoc and correctly tuning them to a **specific environment** might drastically increase performance.

Nikolaus Correll



Program Synthesis Kolmogorov

- Searching program space *exhaustively* is impossible
(Turing: Halting Problem)



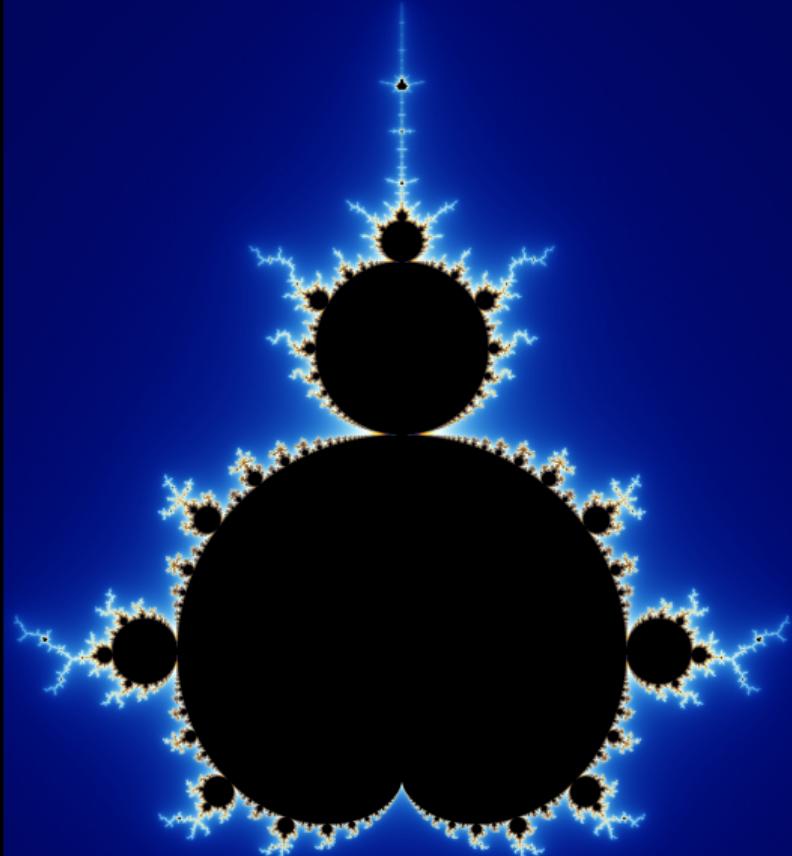
Program Synthesis

Kolmogorov

- Searching program space *exhaustively* is impossible
(Turing: Halting Problem)

Schmidhuber

- Optimal Ordered Solvers



Program Synthesis

Kolmogorov

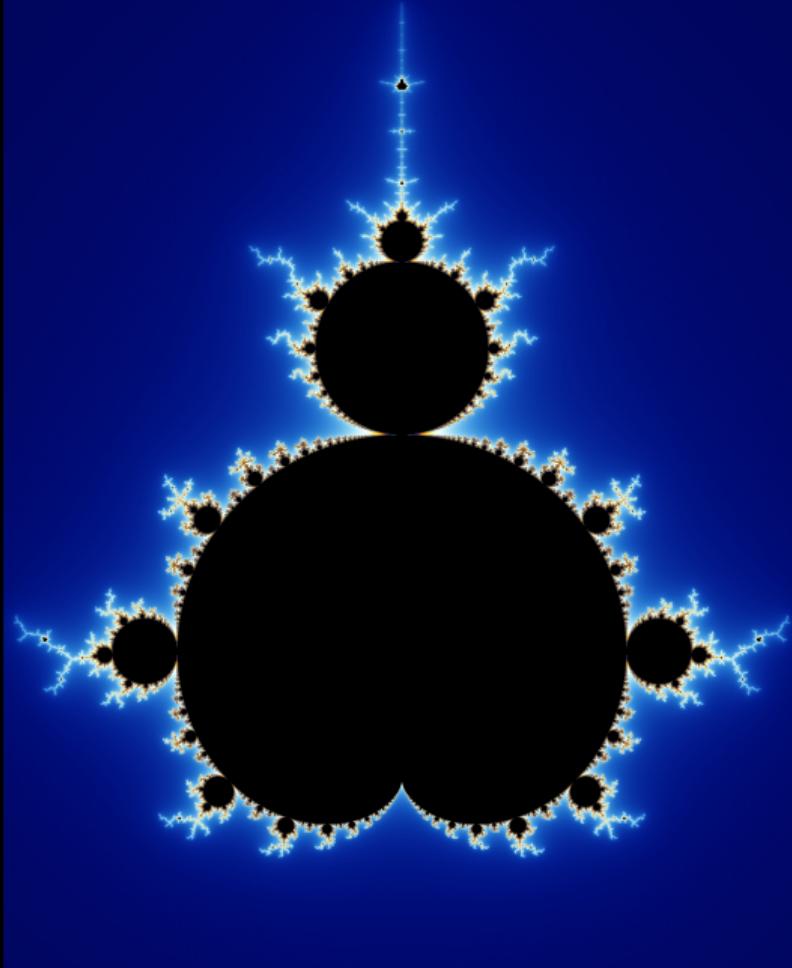
- Searching program space *exhaustively* is impossible
(Turing: Halting Problem)

Schmidhuber

- Optimal Ordered Solvers

Lake & Tenenbaum

- Representation Matters



Many Problems Are
Program Learning In
Disguise

Approach Template Example

Problem
Representation
Optimization



Approach Template Example

Problem
Representation
Optimization

Image Classification



Approach Template Example

Problem

Image Classification

Representation

Network Weights

Optimization



Approach Template Example

Problem

Image Classification

Representation

Network Weights

Optimization

Gradient Methods



“Meta”

Referring to itself or to the conventions of its genre;
self-referential (Wikipedia Definition)

AutoML-Zero: Evolving Machine Learning
Algorithms From Scratch Real et al., 2020

Problem

Representation

Optimization



AutoML-Zero: Evolving Machine Learning Algorithms From Scratch Real et al., 2020

Problem

Meta Image
Classifier

Representation

Optimization



AutoML-Zero: Evolving Machine Learning Algorithms From Scratch Real et al., 2020

Problem Meta Image Classifier

Representation Tensor Machine

Optimization



AutoML-Zero: Evolving Machine Learning Algorithms From Scratch Real et al., 2020

Problem

Meta Image
Classifier

Representation Tensor Machine

Optimization Reg. Evolution

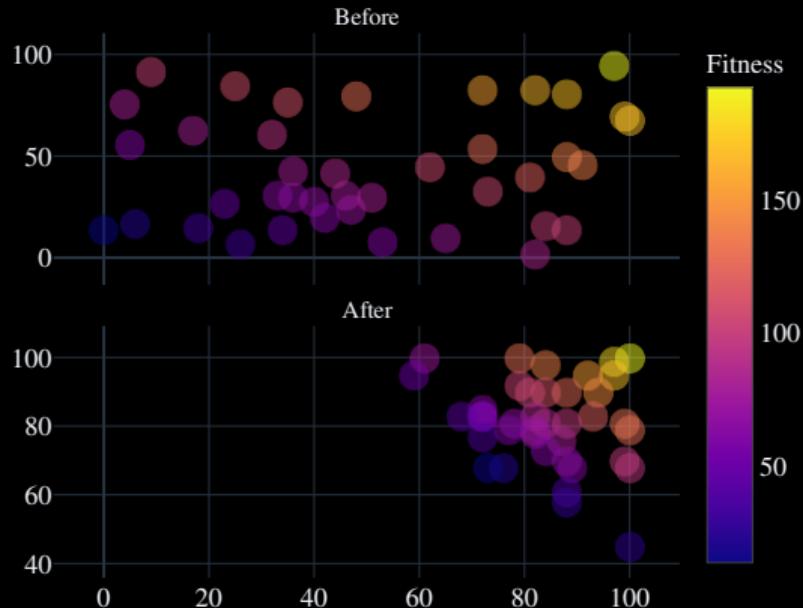


Evolutionary/Genetic Algorithms

Population Optimization

- Selection
- Mutation
- Crossover

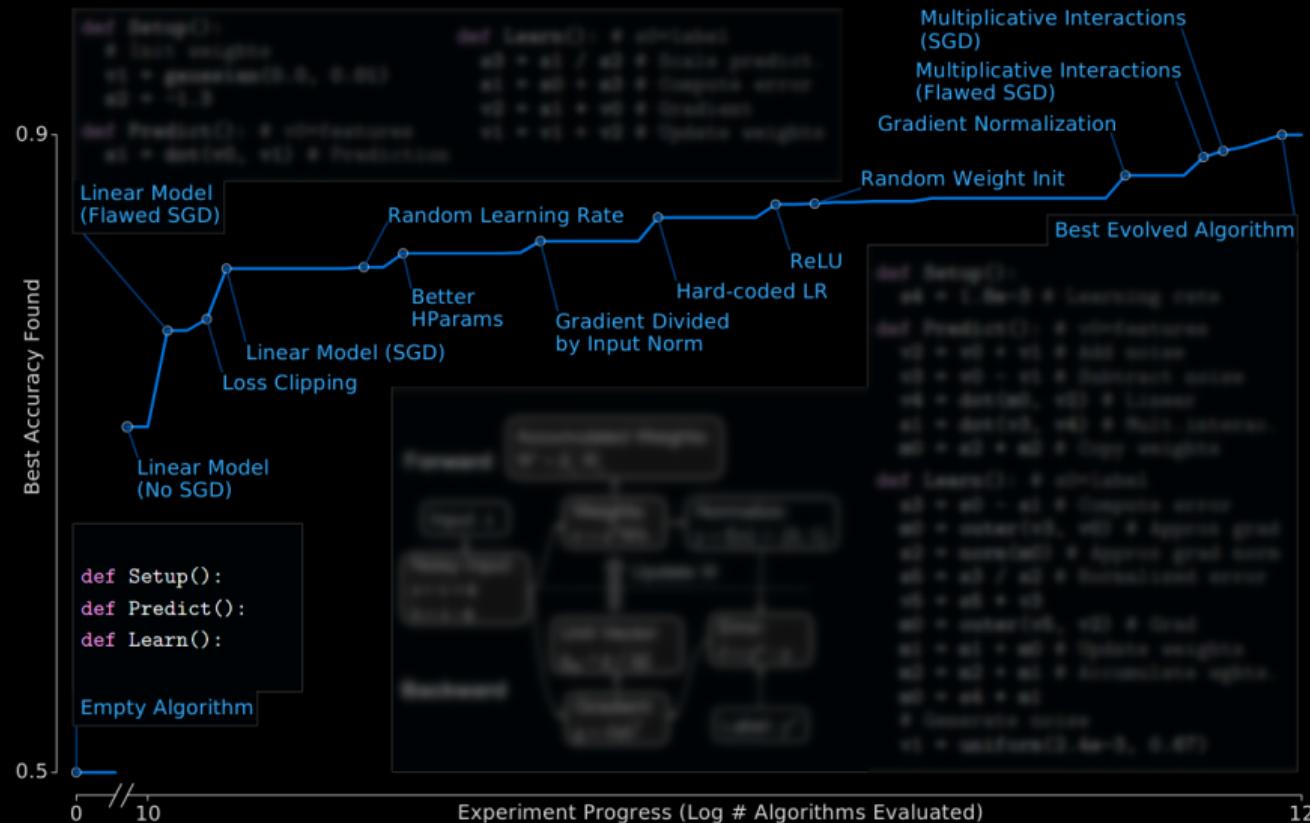
Evolutionary Algorithm Selection



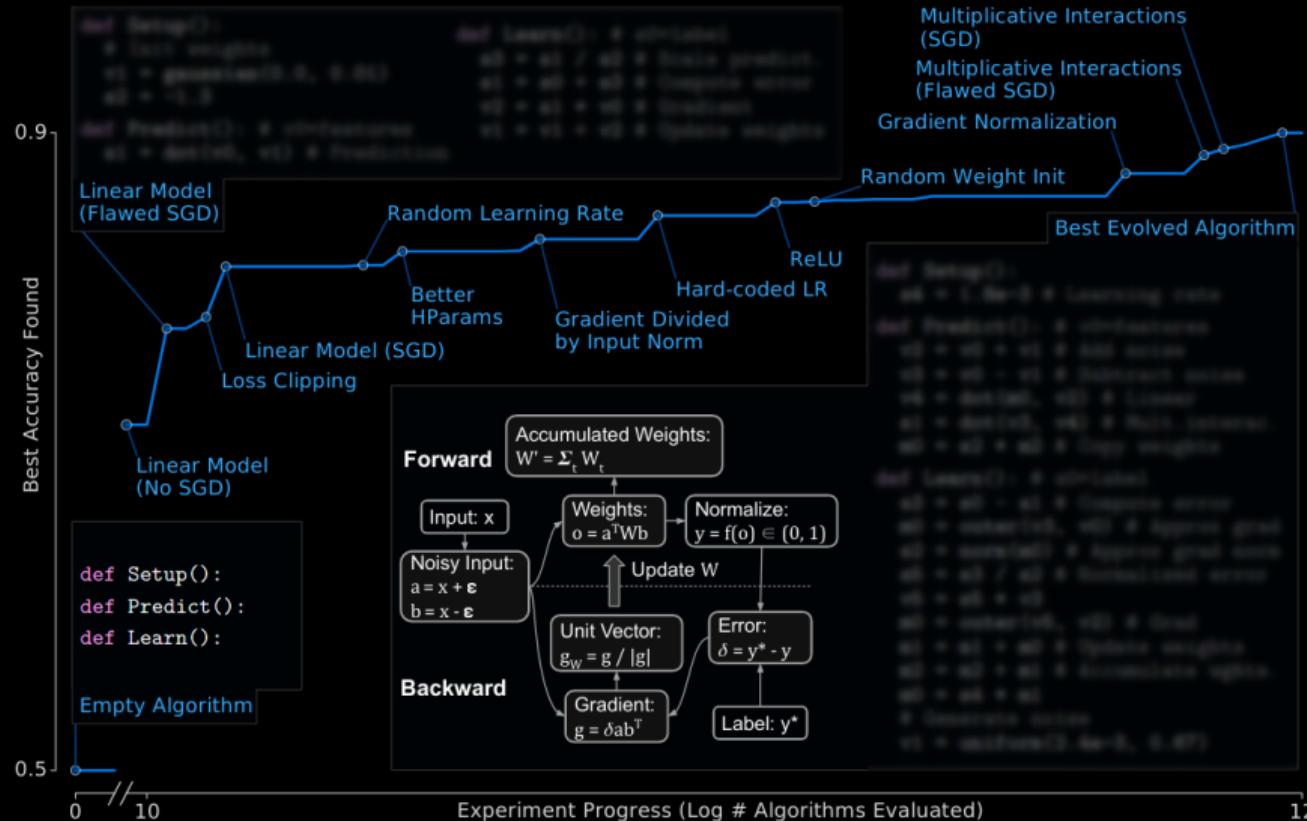
Regularized Evolution Real et al., 2019



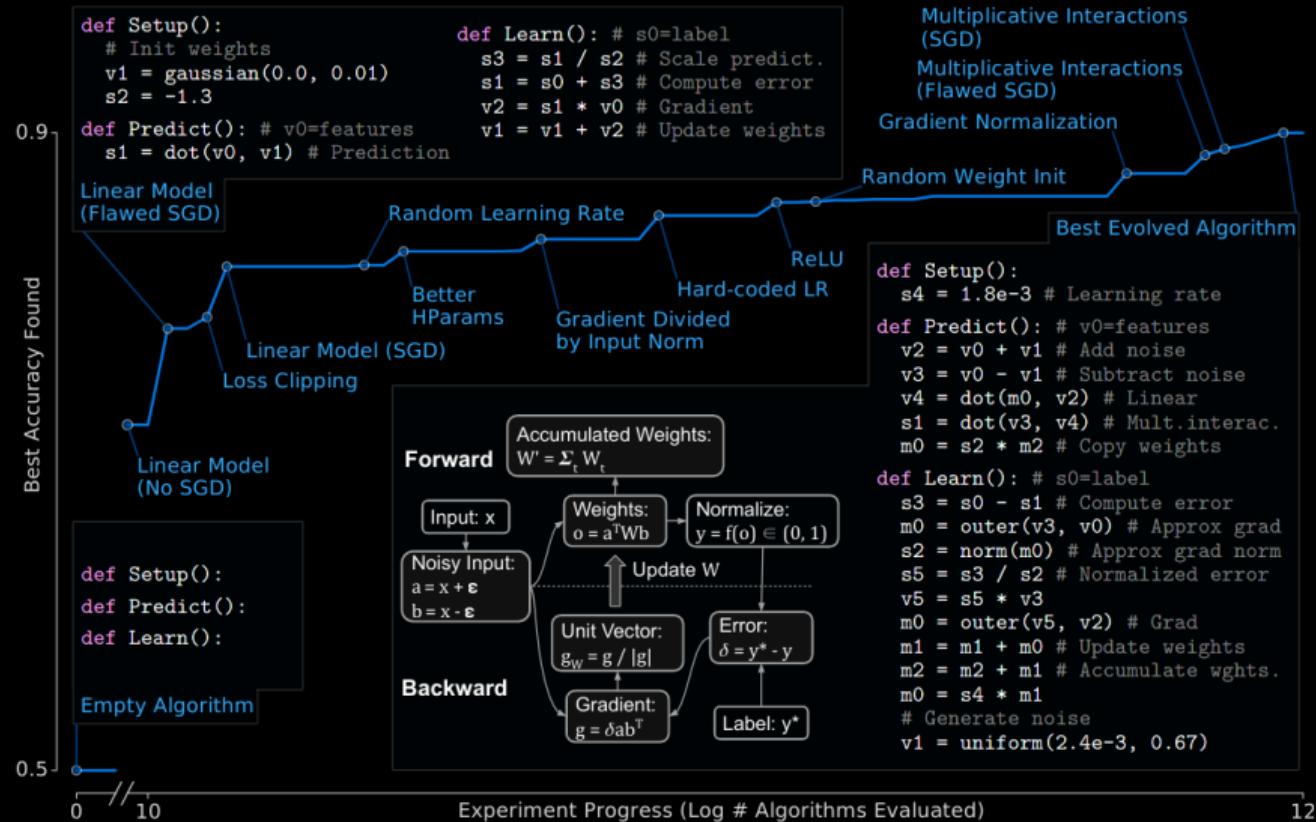
AutoML-Zero Real et al., 2020



AutoML-Zero Real et al., 2020



AutoML-Zero Real et al., 2020



Environment Learning for Path-Planning

Problem Representation Optimization



Environment Learning for Path-Planning

Problem

Path Planning

Representation

Optimization



Environment Learning for Path-Planning

Problem

Path Planning

Representation

Python Code

Optimization



Environment Learning for Path-Planning

Problem

Path Planning

Representation

Python Code

Optimization

Pareto Evolution



Evolvable Python Programs



```
1 # Default RRT*
2 def rrt_star(start, end):
3     K = 1000 # Num nodes
4     tree = Tree()
5     for k in range(K):
6         # Uniformly sample
7         p = sample(bias, end)
8         # Add them to tree
9         tree.extend(p, step)
10        # Rewire nearby paths
11        tree.rewire(p, rad)
12    return tree
```

Evolvable Python Programs

```
1 def build(K=1000):          1 def sample(tree, goal):  
2     step = 50      # evolved 2     # all evolved  
3     rewire = 100 # evolved 3     a = 1/1 # Code surface!  
4     tree = Tree() # const 4     if random() > a:  
5     for k in range(K):    5         return goal  
6         p = sample(tree) 6     else:  
7         tree.extend(p) 7         return (uniform(x),  
8     return tree           8             uniform(y))
```



Evolvable Python Programs

```
1 def build(K=1000):           1 def sample(tree, goal):
2     step = 50    # evolved   2     # all evolved
3     rewire = 100 # evolved   3     a = 1/1 # Code surface!
4     tree = Tree() # const   4     if random() > a:
5     for k in range(K):      5         return goal
6         p = sample(tree)   6     else:
7         tree.extend(p)    7         return (uniform(x),
8     return tree               8             uniform(y))
```

MD5 Hash	α (bias)	r_0 (step)	r_1 (rewire)	K (nodes)
abcde	1.0	50	100	1000
fghij	0.5	100	200	1000

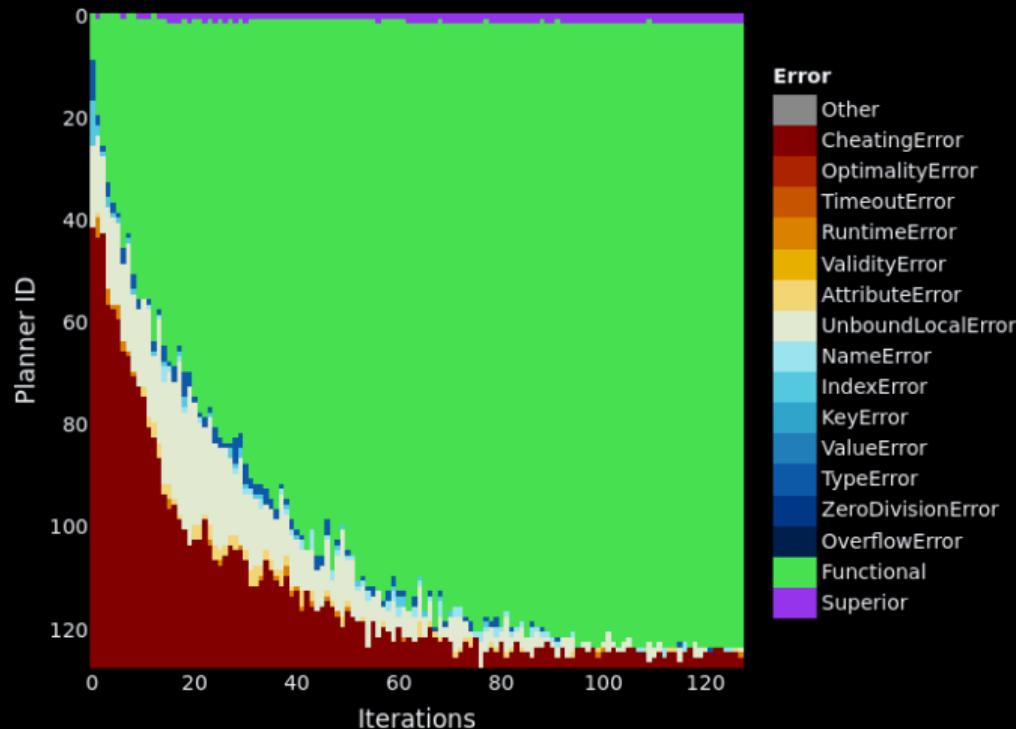


Evolvable Python Grammar

```
1 constant      = -10 .. -1 | 0 | 1 .. 10
2 local         = x           | y           | goal ..
3 funcname     = sin          | cos          | uniform ..
4 binop         = Add          | Sub          | Mult          | Div          | Mod..
5 compop        = Eq           | NotEq        | Lt            | LtE          | Is..
6 unaryop       = UAdd         | USub         | Not           | Inv
7 boolop        = And          | Or
8 expr          = atom         | binary        | unary         | compare
9 stmt          = assign        | return        | logic_stmt   | ast_expr
10 binary        = [atom binop expr]
11 unary         = [unaryop expr]
12 compare       = [atom compops atoms]
13 if            = [expr body body]
```



Evolvable Python Programs

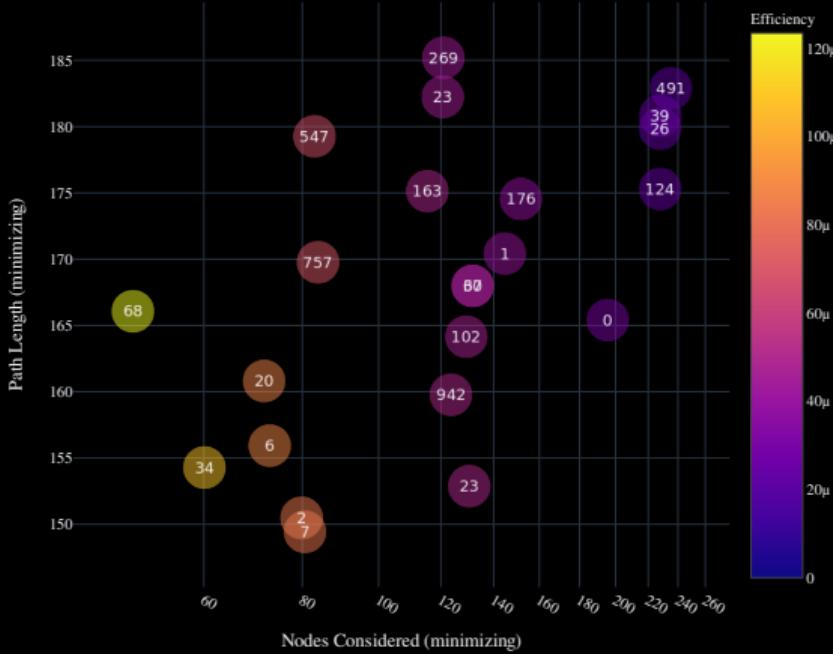


Multiple Objectives

Nodes Path Length

...

Evolved Computer Programs, Labeled By Age



Pareto Fronts

- Vectorized Fitness

$$v_i = \{nodes, path\}$$

- Pareto Equivalence

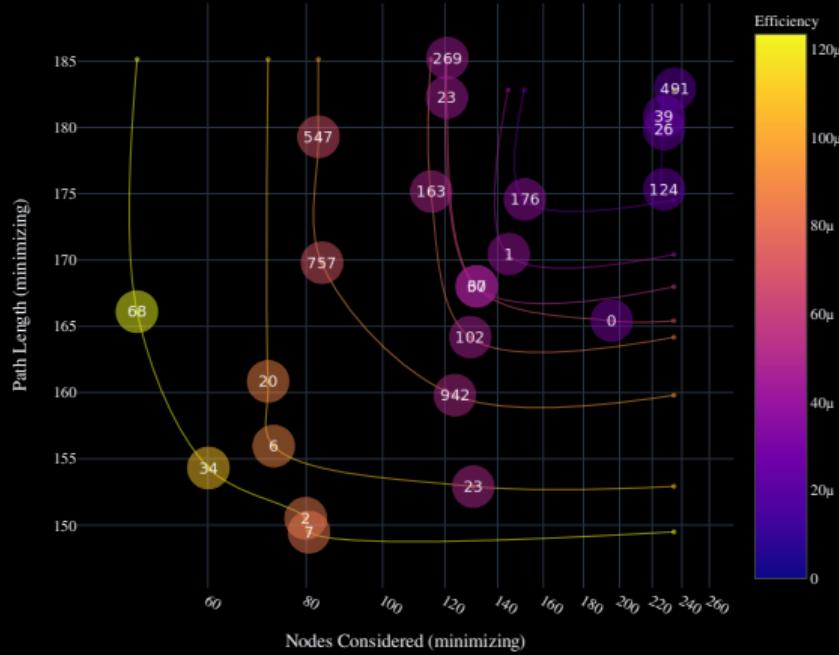
$$v_1 = \{80, 151\}$$

$$v_2 = \{60, 155\}$$

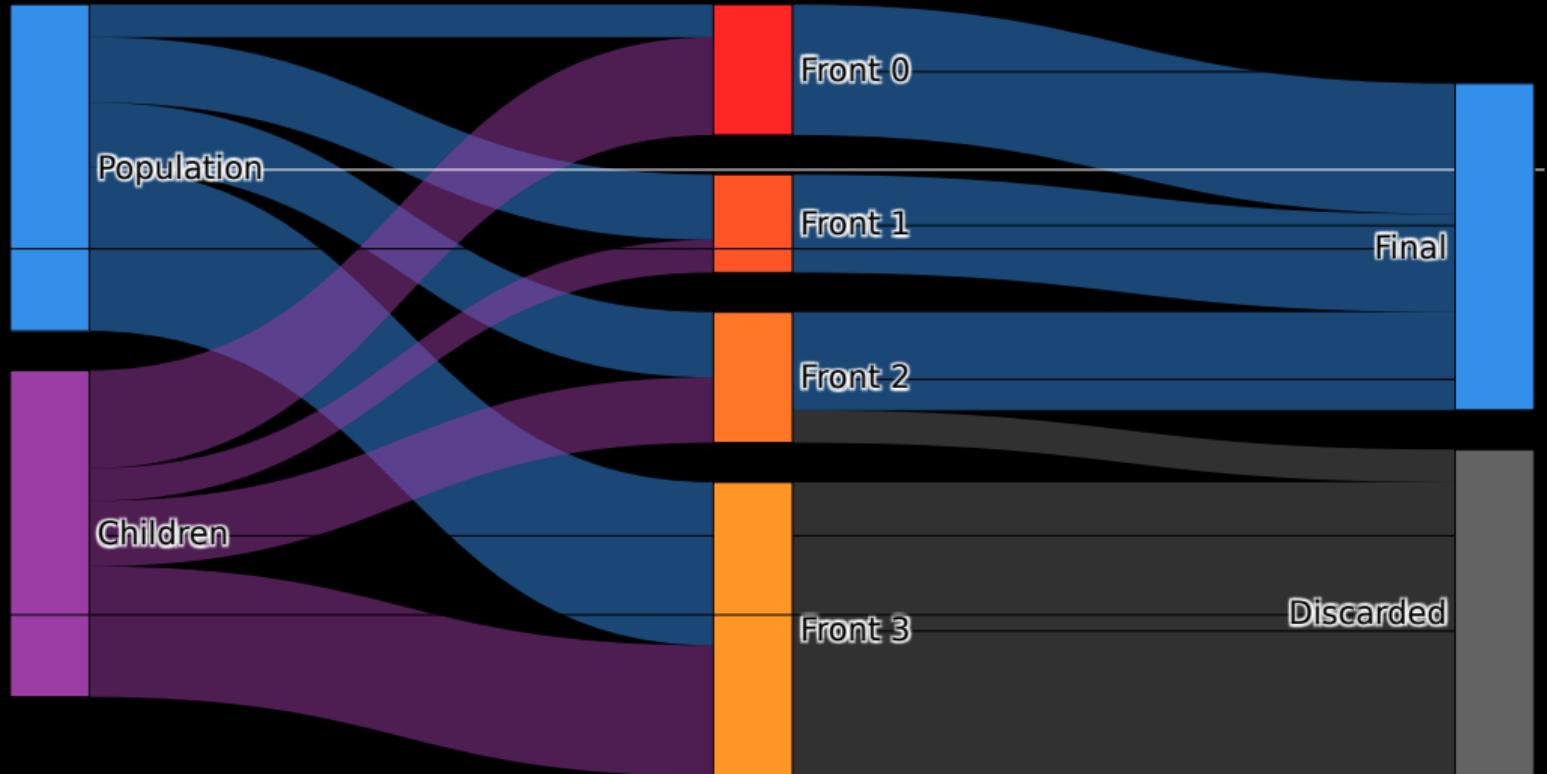
$$v_3 = \{50, 165\}$$

$$v_4 = \{130, 153\}$$

Evolved Computer Programs, Labeled By Age



Pareto Evolution (Srinivas and Deb, 1994)



Pareto Ev. (NSGA2)

- Multi-objective
- Bandwidth: ***n***

Regularized Ev.

- Single-objective
- Bandwidth: **1**



Pareto Ev. (NSGA2)

- Multi-objective
- Bandwidth: n

$$O(g * m * n^2)$$

Regularized Ev.

- Single-objective
- Bandwidth: 1

$$O(g * p)$$

Where g = generations
 m = objectives
 n = total population size
 p = sample size



Pareto Ev. (NSGA2)

- Multi-objective
- Bandwidth: n

$$O(g * m * n^2)$$

$O(m * n)$ per change

Regularized Ev.

- Single-objective
- Bandwidth: 1

$$O(g * p)$$

$O(p)$ per change

Where g = generations
 m = objectives
 n = total population size
 p = sample size



Pareto Ev. (NSGA2)

- Multi-objective
- Bandwidth: n

$$O(g * m * n^2)$$

$O(m * n)$ per change

$O(m * p)$ if sampled

Regularized Ev.

- Single-objective
- Bandwidth: 1

$$O(g * p)$$

$O(p)$ per change

$O(p)$ (sampled)

Where g = generations
 m = objectives
 n = total population size
 p = sample size



Specialization Experiment

Streetmaps: Milan_2_512



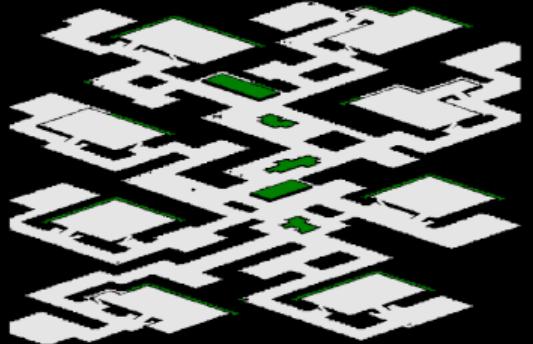
Blank



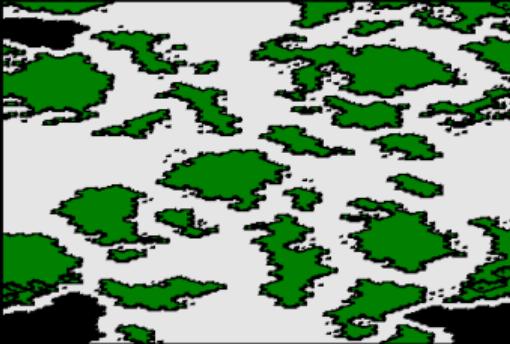
Baldurs Gate: AR0700SR



Starcraft: Enigma



Starcraft: Entanglement

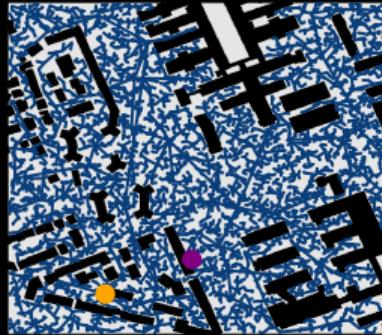


Starcraft: Turbo

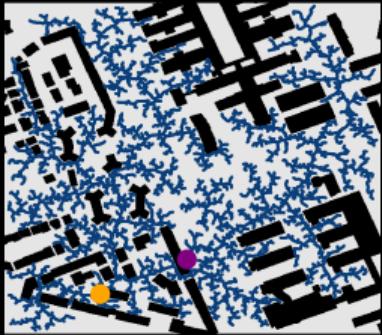


Specialization Results

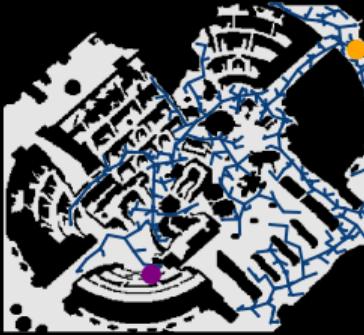
94e7b (Milan best)



a1be6 (Milan worst)



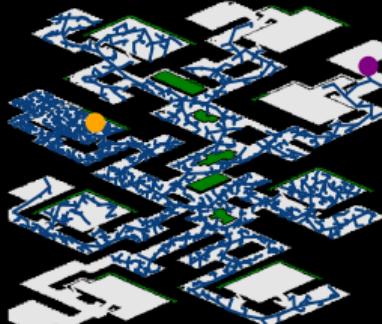
49980 (Validate)



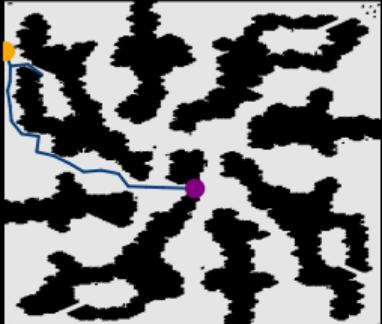
Initial (Validate)



49980 (Enigma best)



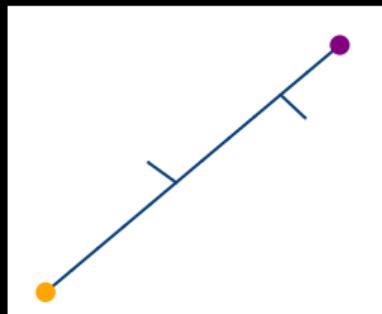
6e799 (Turbo best)



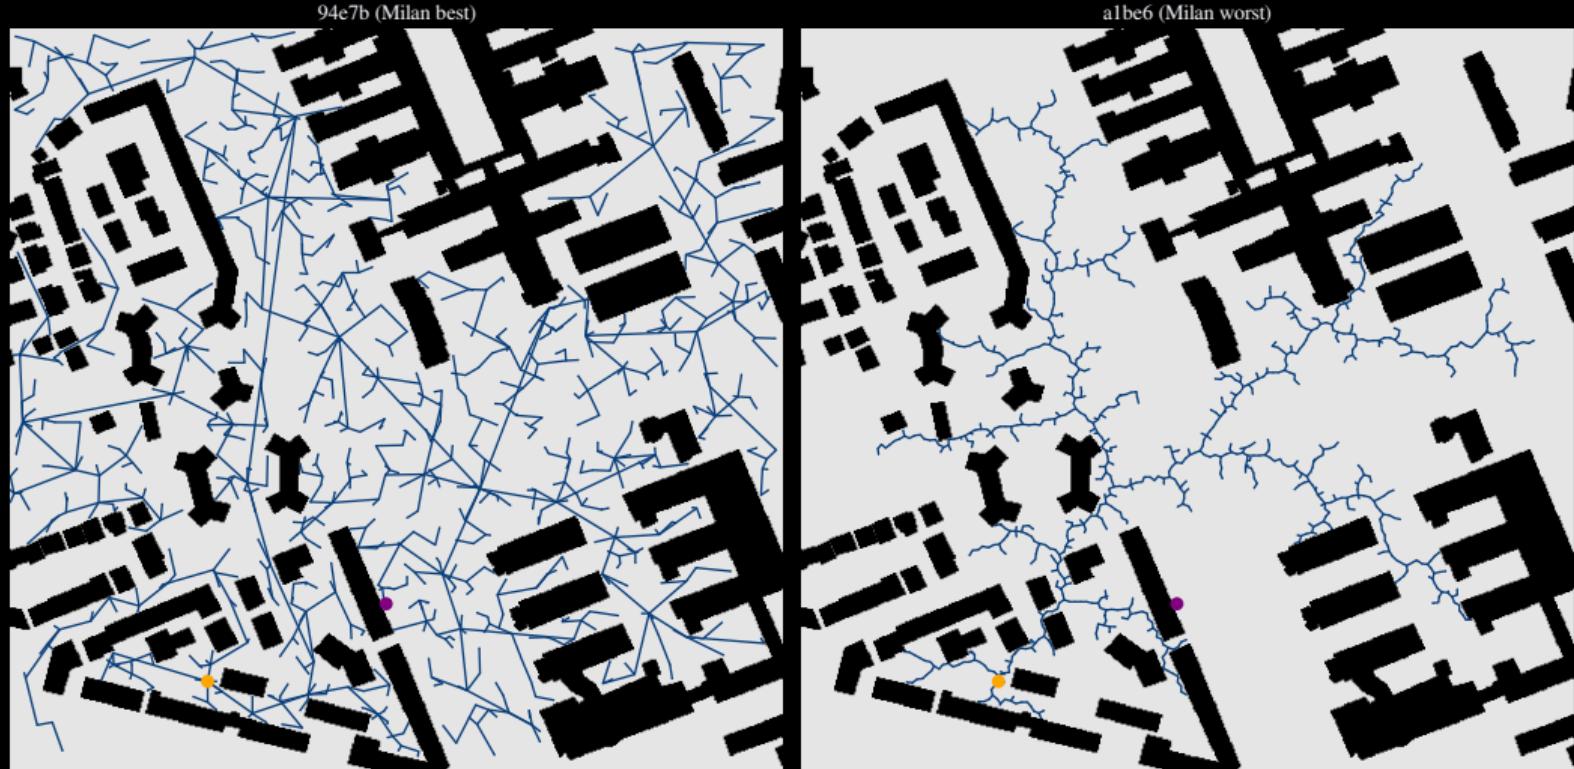
c5b4f (Entanglement best)



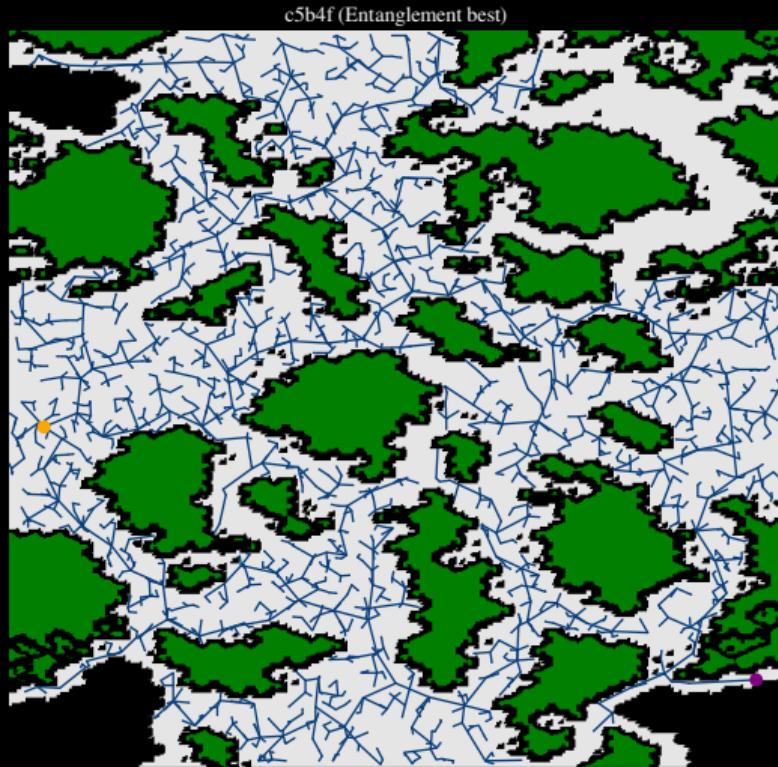
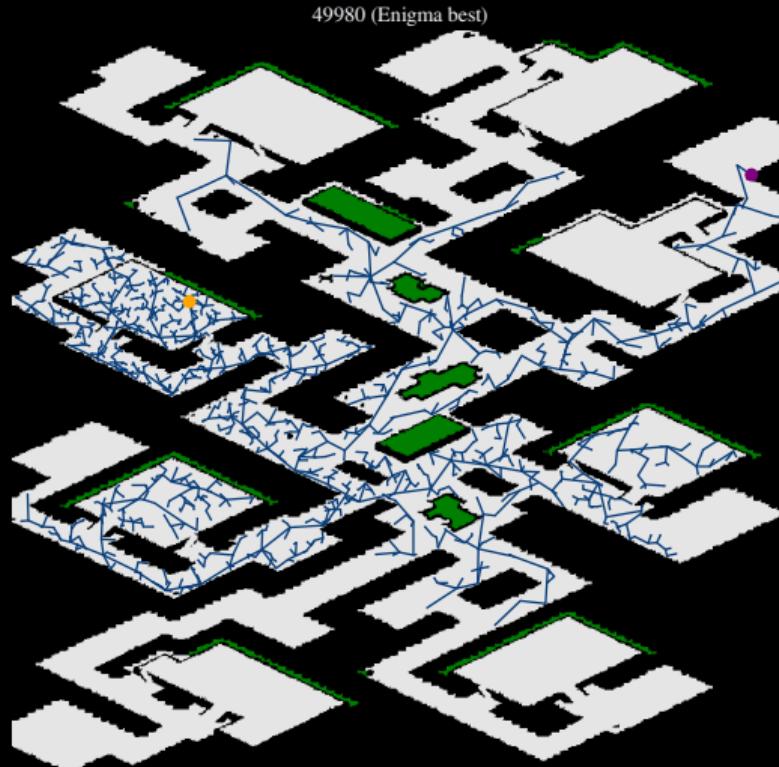
70458 (No Obstacles best)



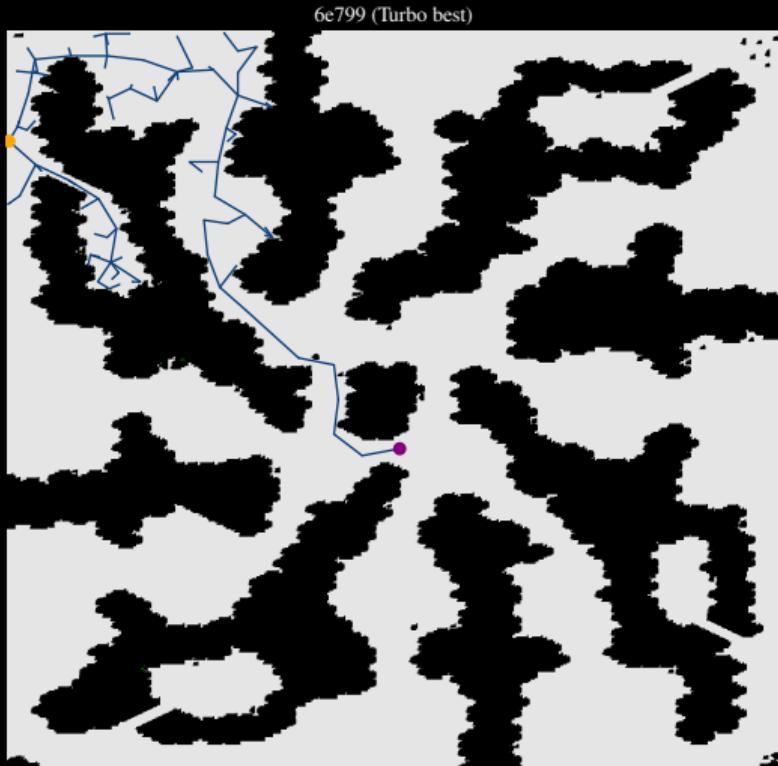
Milan Best/Worst



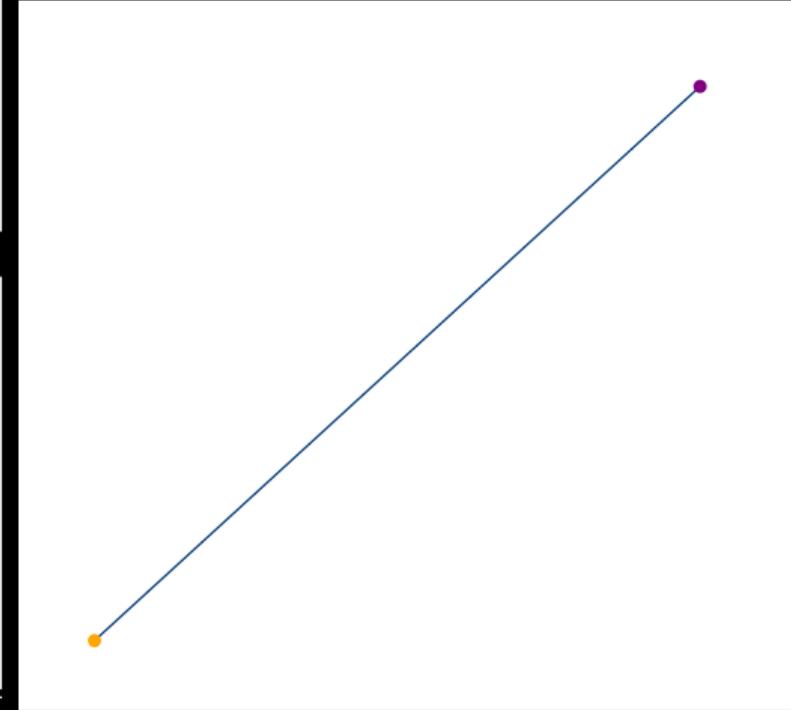
Enigma & Entanglement Best



Turbo and No-Obstacles



70458 (No Obstacles best)

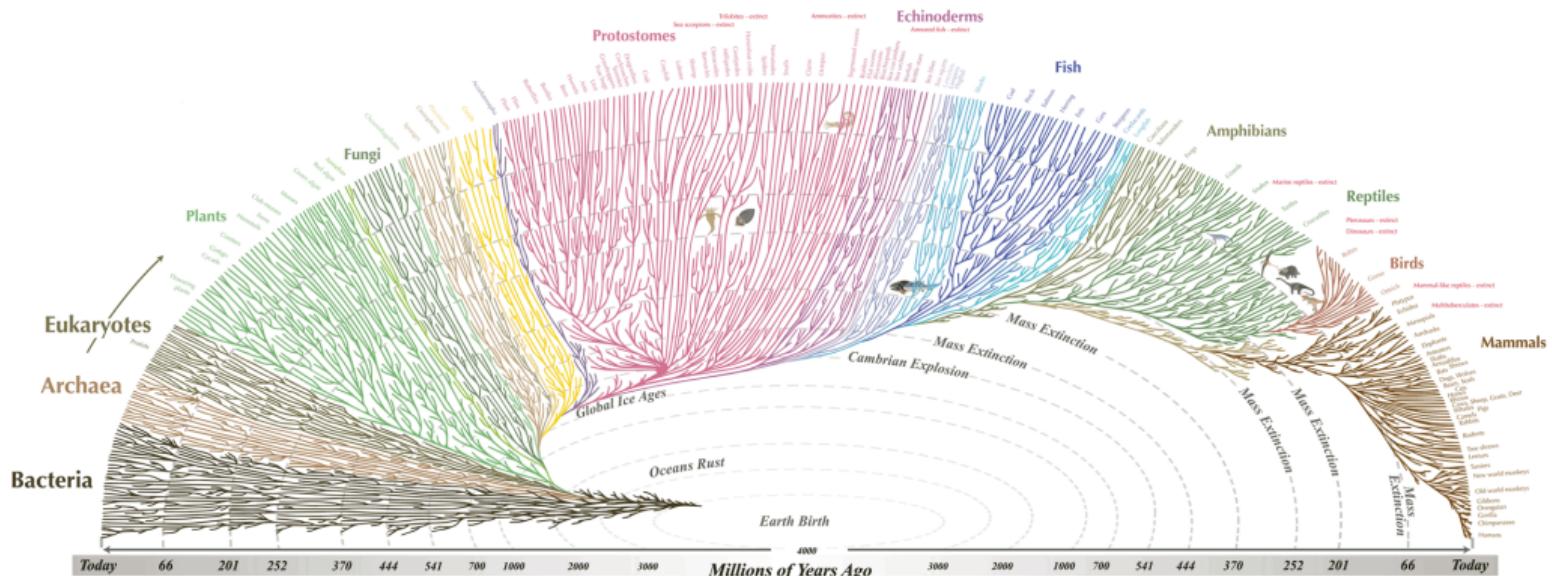


Specialization Results

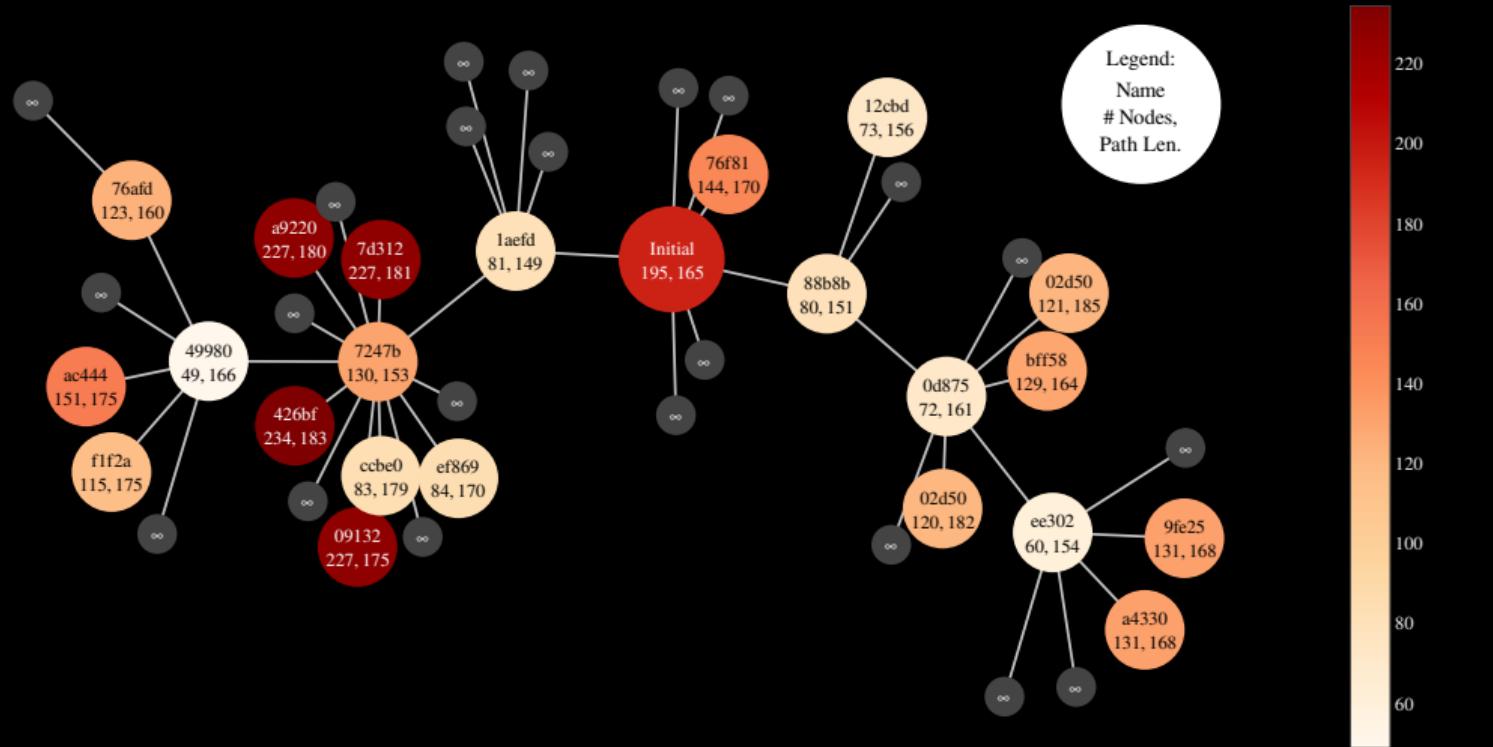
Best-in-class RRT* Statistics (Nodes)

Map	$-\Delta\% \text{ Nodes}$	Edit Depth	Hash	r_0 (step size)	α (bias)
Milan	82.35%	4	94e7b	50,000	7/4 (>1)
Enigma	74.87%	3	49980	100	7/8
Entanglement	85.16%	3	c5b4f	100	7/8
Turbo	64.29%	3	6e799	50	4/9
No Obstacles	89.47%	1	70458	50	1/9
Baseline	0.0%	0	Initial	50	1.0





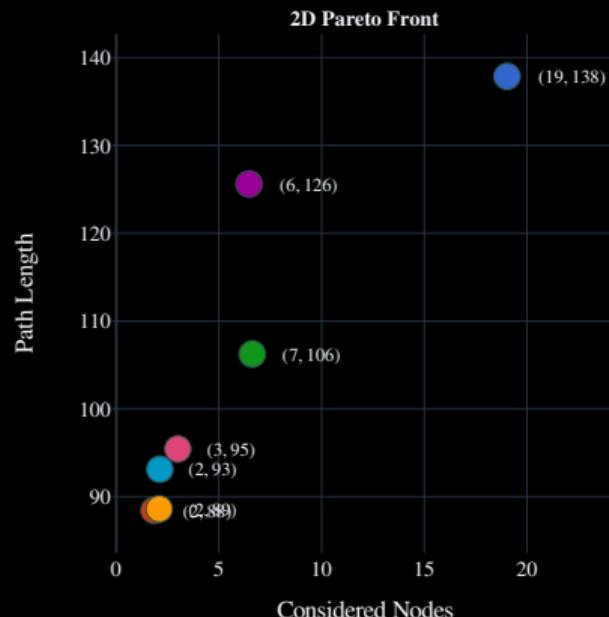
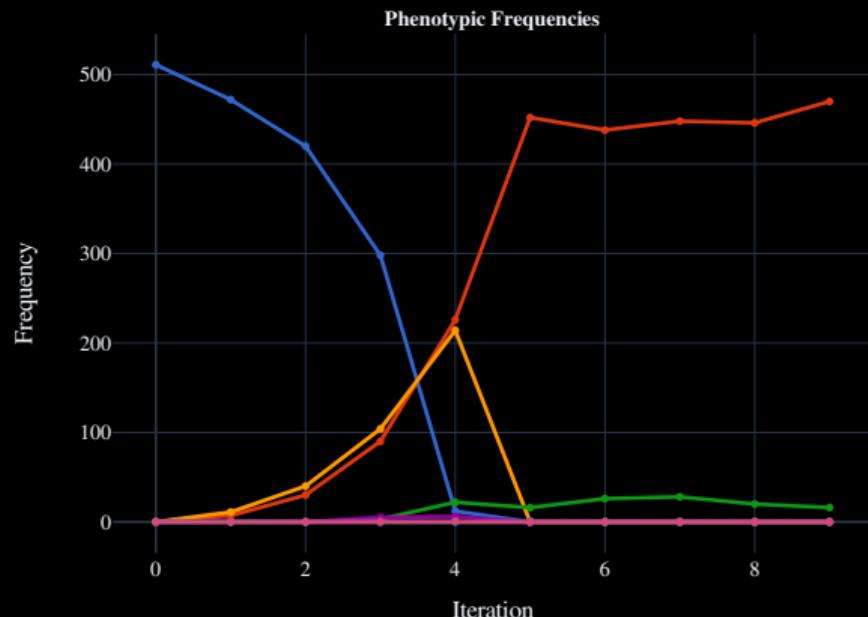
Natural evolution



Evolved Computer Programs

Population Dynamics: Blank Map

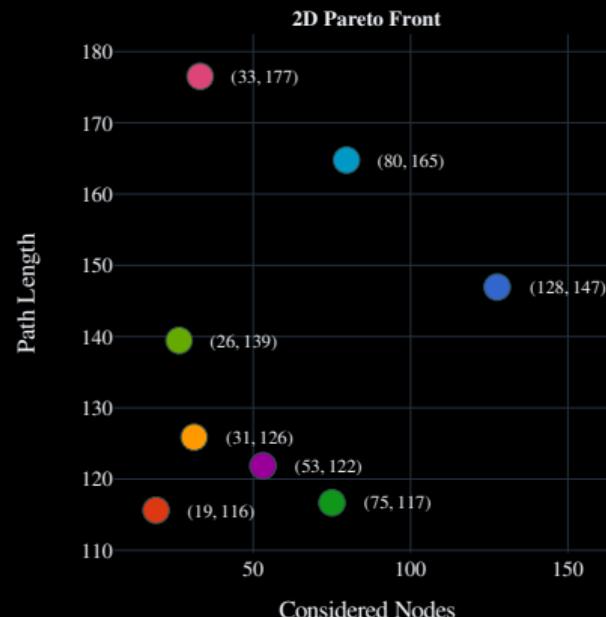
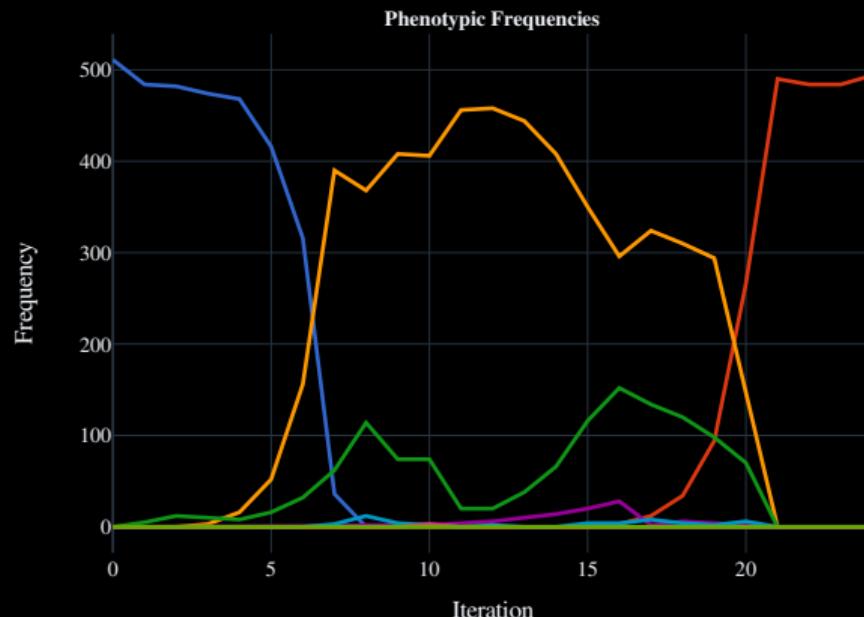
Program Fingerprints ● Initial ● 70458 ● 9d7e1 ● ee05a ● 67a69 ● d7379 ● 32ca3



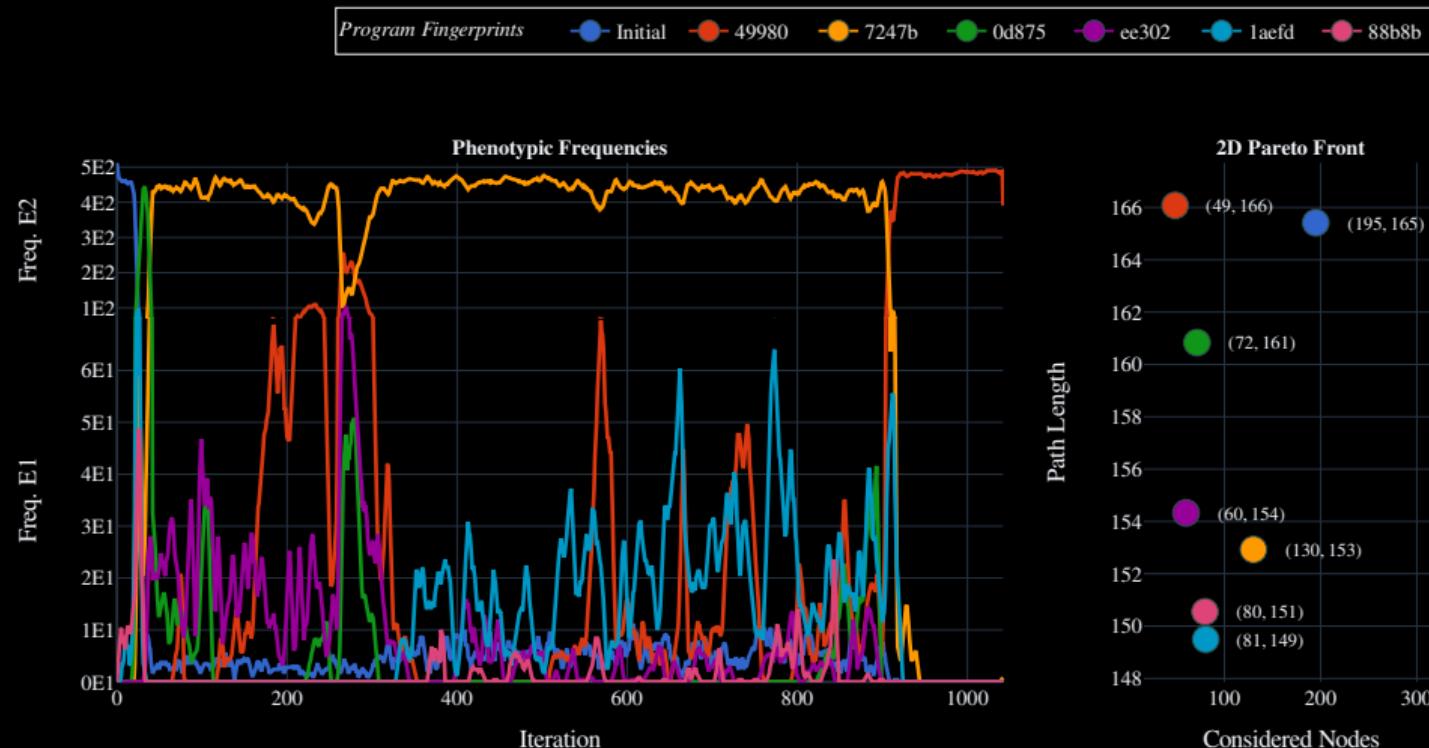
Population Dynamics: Entanglement

Program Fingerprints

Initial (Blue), c5b4f (Red), 7b368 (Orange), 9859b (Green), 0309e (Purple), 48f1b (Cyan), c4566 (Pink), 398f4 (Light Green)



Population Dynamics: Starcraft Enigma

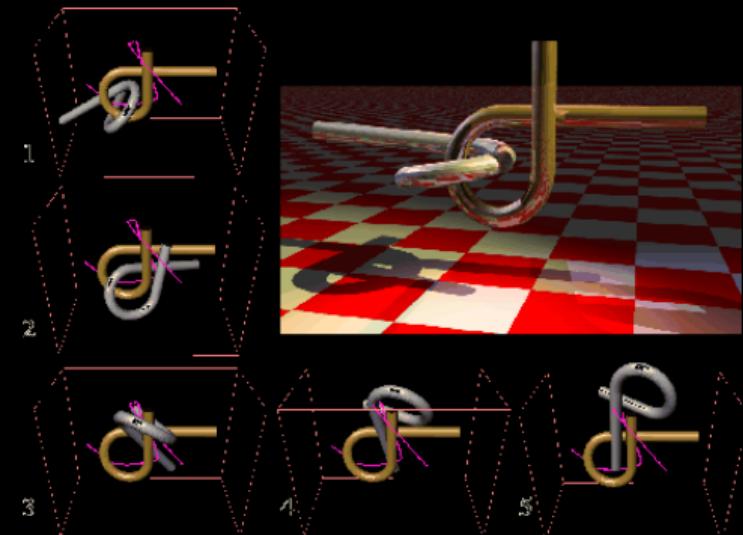


Future Work

Algorithms

Environments

Optimization



LaValle, 2006

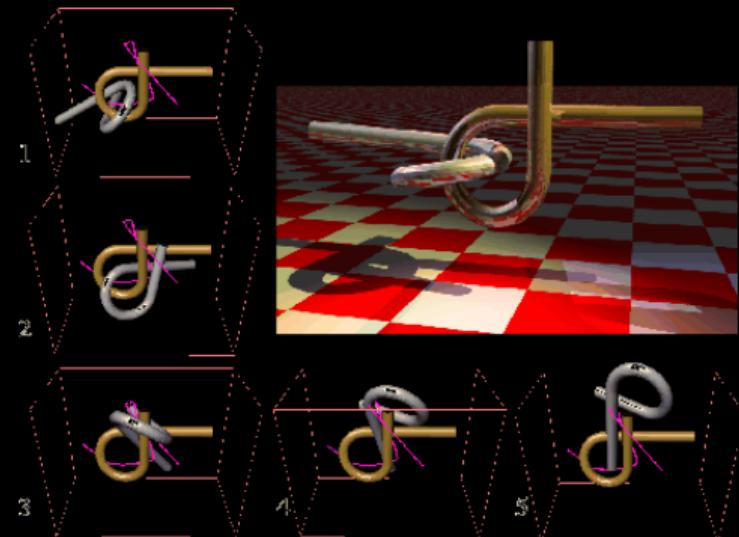
Future Work

Algorithms

- Advanced Route Planning
- Advanced Sampling

Environments

Optimization



LaValle, 2006

Future Work

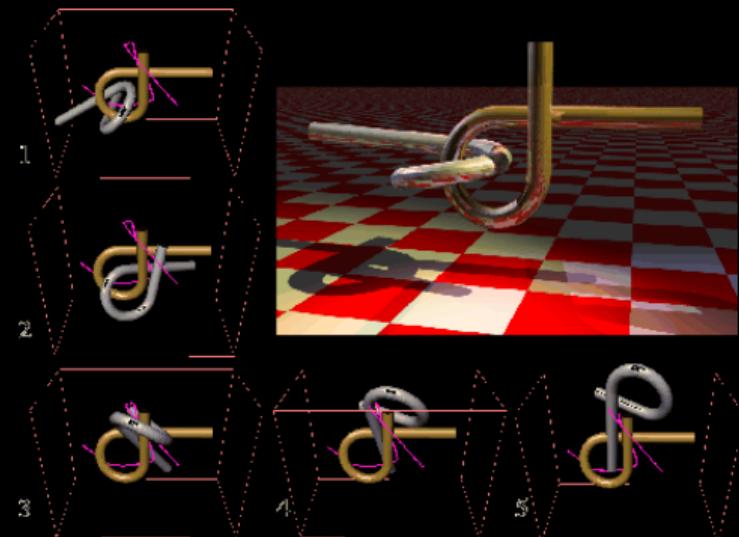
Algorithms

- Advanced Route Planning
- Advanced Sampling

Environments

- High-Dimensions

Optimization



LaValle, 2006

Future Work

Algorithms

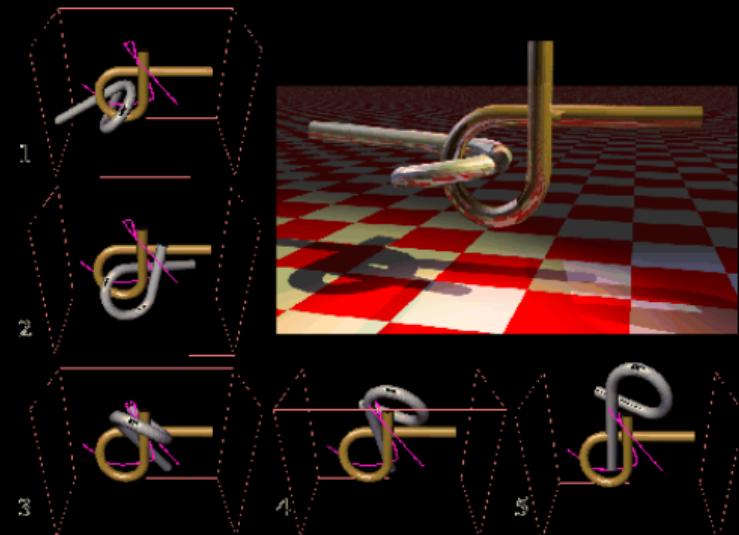
- Advanced Route Planning
- Advanced Sampling

Environments

- High-Dimensions

Optimization

- Alt. Pareto Evolution



LaValle, 2006

Program Learning Template

Program Learning Problems

Problem	Representation	Optimization
Image Classification	NN Weights	Gradient
Char. Recognition	Probabilistic Program	Bayes. Inference
Arch. Search	Tensor Register Machine	Regularized Evolution
Path Planning	Python Program	Pareto Evolution

Wait, it's all
program
learning?



Always has been.

No Free Lunch

No Free Lunch

Meta learning and specialization

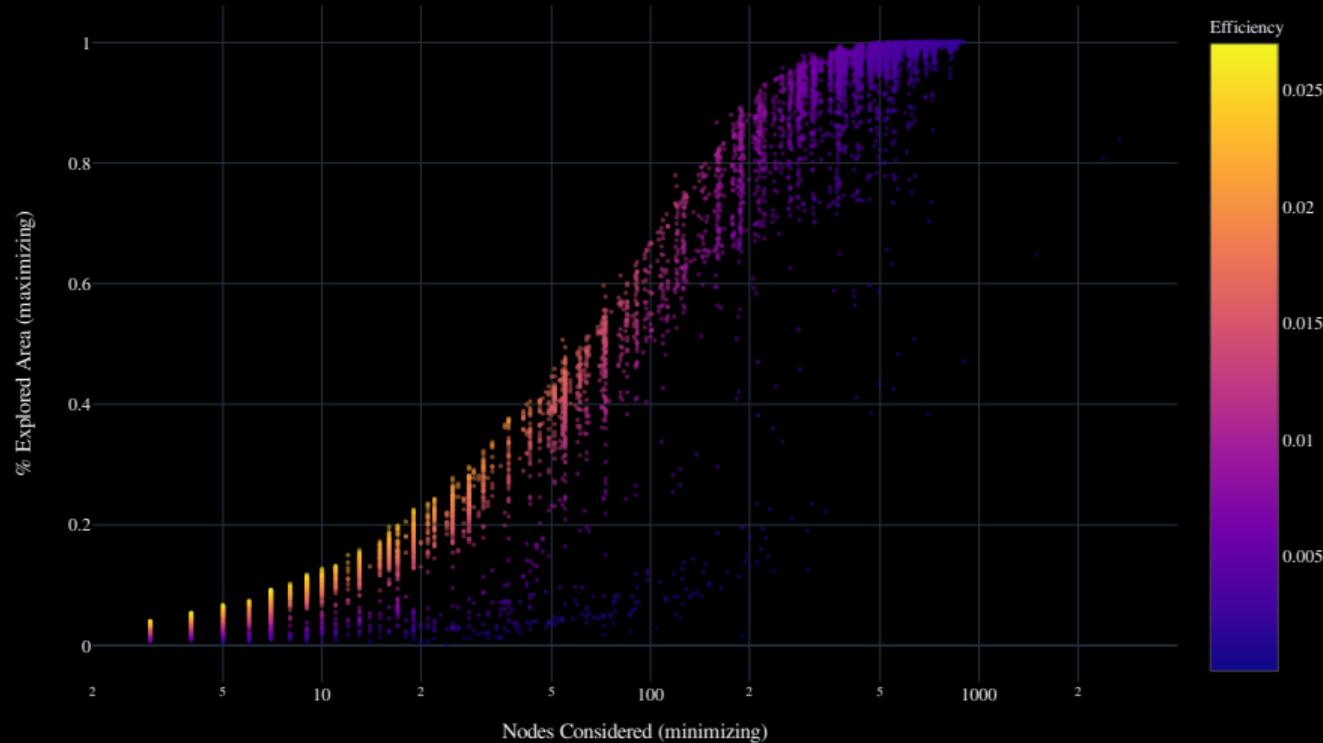
Questions?

Wait, it's all
program
learning?



Always has been.

Efficient Exploration



Baldur's Gate Validation

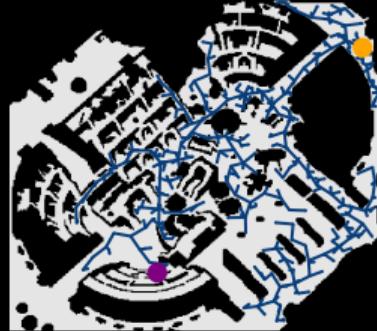
94e7b (Milan best)



a1be6 (Milan worst)



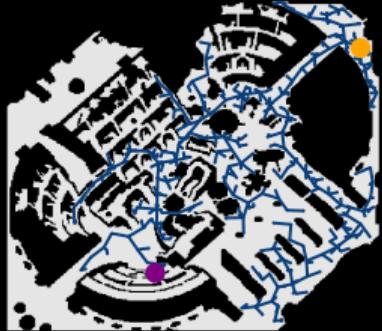
49980 (Convergent)



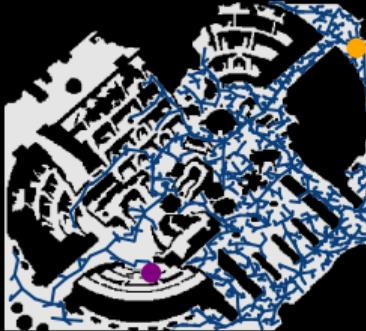
Initial (Validate)



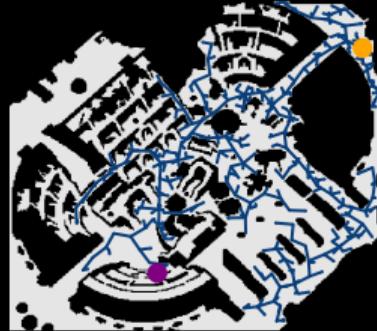
49980 (Enigma best)



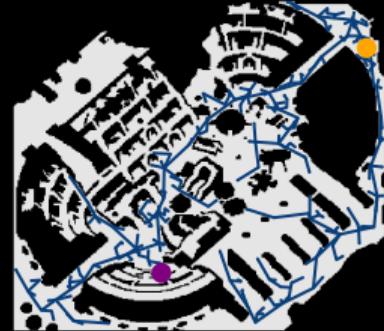
6e799 (Turbo best)



c5b4f (Entanglement best)

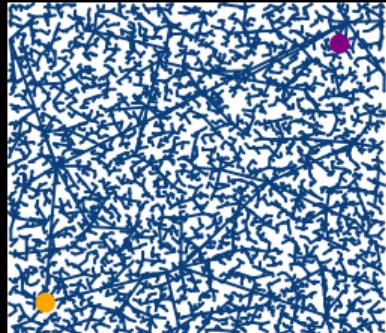


70458 (No Obstacles best)

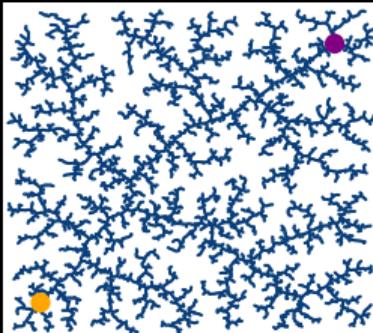


Blank Validation

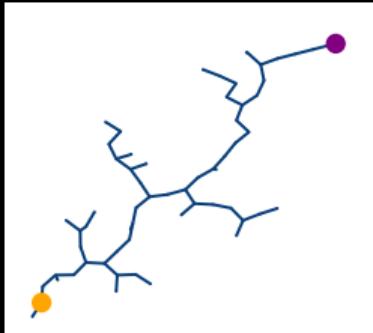
94e7b (Milan best)



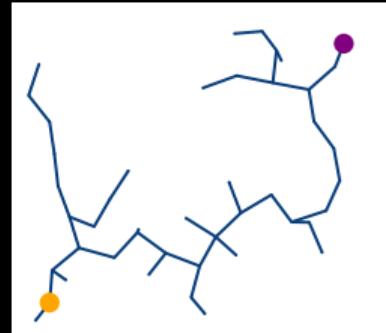
a1be6 (Milan worst)



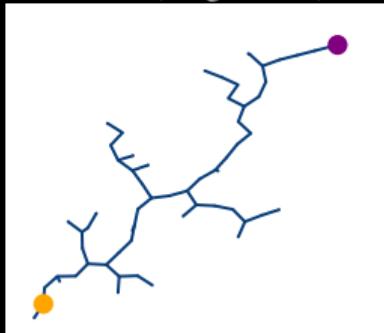
49980 (Convergent)



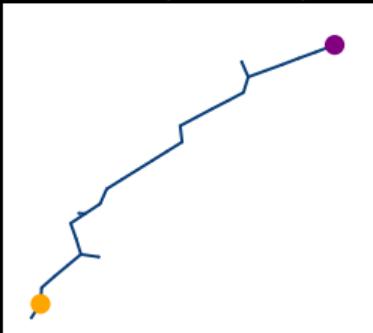
Initial (Validate)



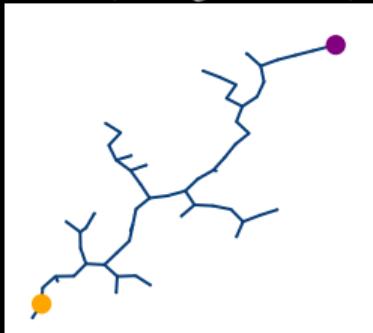
49980 (Enigma best)



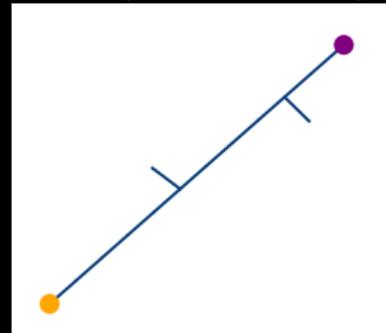
6e799 (Turbo best)



c5b4f (Entanglement best)



70458 (No Obstacles best)



- Bast, H., Delling, D., Goldberg, A., Müller-Hannemann, M., Pajor, T., Sanders, P., Wagner, D., & Werneck, R. F. (2016). Route planning in transportation networks. *Algorithm engineering* (pp. 19–80). Springer.
- Kavraki, L. E., Svestka, P., Latombe, J.-C., & Overmars, M. H. (1996). Probabilistic roadmaps for path planning in high-dimensional configuration spaces. *IEEE transactions on Robotics and Automation*, 12(4), 566–580.
- LaValle, S. M. (2006). *Planning algorithms*. Cambridge university press.
- LaValle, S. M., & Kuffner Jr, J. J. (2001). Randomized kinodynamic planning. *The international journal of robotics research*, 20(5), 378–400.
- Real, E., Aggarwal, A., Huang, Y., & Le, Q. V. (2019). Regularized evolution for image classifier architecture search. *Proceedings of the aaai conference on artificial intelligence*, 33(01), 4780–4789.
- Real, E., Liang, C., So, D., & Le, Q. (2020). Automl-zero: Evolving machine learning algorithms from scratch. *International Conference on Machine Learning*, 8007–8019.
- Srinivas, N., & Deb, K. (1994). Multiobjective optimization using nondominated sorting in genetic algorithms. *Evolutionary computation*, 2(3), 221–248.